


Diagnostics

UBCO MDS — DATA 570



Potential Problems



- When we fit a linear regression model to a particular data set, a number of problems may occur.
- Today we will go through the most common among these.

Non-linearity of the Data



The SLM assumes that the relationship between the predictors and the response fall (\approx) on a straight line.

When the true relationship deviates markedly from linear, then all conclusions/predictions derived from this fit are dubious at best.

Non-linearity of the Data



Residual plots are a useful graphical tool for identifying non-linearity.

SLR: plot the residuals, $e_i = y_i - \hat{y}_i$, versus the predictor x_i .

MLR: plot the residuals versus the predicted (or fitted) values \hat{y}_i .

Non-linearity of the Data



- Ideally, the residual plot will show no discernible pattern.
- The presence of a pattern may indicate a problem with some aspect of the linear model.
- A “good” plot should produce randomly dispersed residuals around the horizontal axis





Figure: “Bad” (left) and “good” (right) residual plots.

Residual Plots



- To help us identify any trends, a red smooth line is fit to the residuals
- In the ideal scenario, this smooth line should be a horizontal line at 0.
- In presence of non-linear associations, we could try and fit a model with *transformed* predictors, eg. $\log(X)$, \sqrt{X} , and X^2 , (as discussed in Lecture 5)

Correlation of error terms




- As discussed in Lecture 2, the linear regression model assumes that the error terms, e_1, e_2, \dots, e_n , are independent.
- If error terms are correlated, then the estimated standard errors will tend to underestimate the true standard errors.
- In other words, if the error terms are correlated, we may have an unwarranted sense of confidence in our model.

Correlation of error terms



- Correlated errors can occur when multiple measurements are taken on the same subject over time (i.e. *time series* data)
- To investigate this assumption violation, we plot the residuals from our model as a function of time.
- Other examples can be when our data contain
 - members of the same family
 - people exposed to the same environmental factors



Non-constant variance





- Non-constant variance of error terms, i.e. $Var(e_i) = \sigma^2$, is another important assumption of the linear regression model.
- This assumption is often violated, eg. the variances of the error terms may increase with the value of the response.

Non-constant variance



- Non-constant variance, i.e. *heteroscedasticity* can be identified through the residual plot.
- Often this violation takes shape as a funnel/fan shape in the residual plot.
- In presence of non-constant variance, we could try transforming the *response* Y using a concave function eg. $\log(Y)$, \sqrt{Y} .



Outliers



- An outlier is a point for which y_i is far from the predicted value \hat{y}_i (see next slide for example)
- Outliers can be genuine or be a result of, for example, incorrect reading/recording.
- If an outlier is believed to be a result of an error, then one might simply remove that observation.
- However, genuine outliers may indicate a deficiency with the model, eg. a missing predictor.




Figure: The red/blue line shows the fit with the outlier included/excluded.

Outliers



- As depicted in the previous slide, it can happen that outliers do not greatly affect on the least squares regression line.
- It can however affect things like RSE, p -values, and R^2 values.
- Again, residual plots can be useful in identifying outliers.

Outliers



- To avoid arbitrary cut-offs in a residual plot, sometimes we turn to the so-called *studentized residuals*
- These scaled residuals are produced by dividing each residual e_i by its estimated standard studentized error, $se(e_i) = \sqrt{MSE(1 - h_i)}$
- Observations whose studentized residuals are greater than 3 are then flagged as potential outliers.

High-leverage points



- While outliers produce unusual y_i values, observations with **high leverage** have unusual x_i values.
- The inclusion/exclusion of high leverage points tend to have a higher impact on estimated regression line
- It is very important to identify such points as they have the potential to invalidate the entire fit.

High-leverage points





Figure: Left: example of outlier Right: example of high-leverage. The red (blue) line shows the fit with the outlier/high-leverage included (excluded)





Figure: Plotting the studentized residuals versus h_i flags observation 41 as an outlier as well as a high-leverage point, while observation 20 is an outlier with low leverage.

- For SLR high leverage can be easily identified by looking for observations having predicted values outside of the “normal” range
- With MLR, these points can be harder to identify.



High-leverage observations will have larger values of the so-called *leverage statistic*. For SLR:

$$h_i = \frac{1}{n} + \frac{x_i - \bar{x}}{\sum_j^n (x_j - \bar{x})^2} \quad (1)$$

As the distance between x_i and \bar{x} increase, so to does the h_i .

For MLR, leverage is given the i th diagonal element of the *hat* matrix:

$$H = X(X^T \cdot X)^{-1}X^T \quad (2)$$

In words: how far the vector $(X_{i1}, X_{i2} \dots X_{ip})$ is from $(\bar{X}_1, \bar{X}_2 \dots \bar{X}_p)$, with distance measured in standard deviation units.

A general guideline is to use $h_i > 2(p + 1)/n$ as an indicator for high leverage (note: $1/n \leq h_i \leq 1$)

Influential points



- Observations having a relatively large effect on the regression model's predictions are called **influential** observations
- A high leverage point is not necessarily an influential point.
- An influential point is not necessarily a high leverage point.

Influential points




Figure: Left: Influential high leverage outlier Right: High leverage

Influential points



- A typical measure of influence is the Cook's D-statistic.
- The Cook's distance statistics, for observation i :

$$D_i = \frac{e_i^2}{MSE \cdot d} \left[\frac{h_i}{(1 - h_i)^2} \right] \quad (3)$$

where d is the dimension of your data (ie. $X_{n \times d}$)

- Influential observations will have high Cook's distance score.

Collinearity



- **Collinearity** occurs when two or more predictor variables are closely related to one another.
- Simple bivariate scatterplots can show us correlations between predictors
- However, one variable may be correlated with some *linear combination* of two or more other variables (**multicollinearity**)

Collinearity



- Collinearity reduces the accuracy of the estimates of the regression coefficients, ie. increase $SE(\hat{\beta}_j)$.
- Consequently, the *power* of the hypothesis test—the probability of correctly power detecting a non-zero coefficient—is reduced by collinearity




Figure: ISLR Figure 3.14: Scatterplots of the observations from the **Credit** data set. Left: A plot of **age** versus **limit**. These two variables are not collinear. Right: A plot of **rating** versus **limit**. There is high collinearity

		Coefficient	Std. error	t-statistic	p-value
Model 1	Intercept	-173.411	43.828	-3.957	< 0.0001
	age	-2.292	0.672	-3.407	0.0007
	limit	0.173	0.005	34.496	< 0.0001
Model 2	Intercept	-377.537	45.254	-8.343	< 0.0001
	rating	2.202	0.952	2.312	0.0213
	limit	0.025	0.064	0.384	0.7012

TABLE 3.11. *The results for two multiple regression models involving the Credit data set are shown. Model 1 is a regression of balance on age and limit, and Model 2 a regression of balance on rating and limit. The standard error of $\hat{\beta}_{\text{limit}}$ increases 12-fold in the second regression, due to collinearity.*

Collinearity



- The most straightforward measure of collinearity is called the **Variance Inflation Factor (VIF)**.
- VIF measures the ratio of the variance of $\hat{\beta}_j$ in the full model and the variance of $\hat{\beta}_j$ if it were fit on its own (i.e SLR).
- A $VIF=1$ (smallest possible value) indicates the complete absence of collinearity

Collinearity



- The VIF is calculated as follows:

$$VIF(\hat{\beta}_j) = \frac{1}{1 - R_j^2}$$

where R_j^2 is the R^2 for the regression of X_j on all the other X_s .

- If R_j^2 is close to one, then collinearity is present, and so the VIF will be large.
- Typically values exceeding 10 indicate a problem.

Collinearity



- In practice small amount of collinearity among the predictors is expected
- VIFs do not tell how many collinearities there are, or which variables are included in them.
- There are other more sophisticated measures of collinearity (eg. based on eigenvalues and eigenvectors of the matrix of X s) but those fall outside the scope of this module.

Collinearity



In the face of collinearity, we may decide to:

- eliminate one of the problematic variables our model
- combine the collinear variables together into a single predictor

Diagnostic Plots



If we apply the `plot()` function to the output from a `lm()` (see `?plot.lm` for details), four diagnostic plots are produced:

1. Residuals vs Fitted
2. Normal Q-Q plot
3. Scale-Location
4. Residuals vs Leverage

Normal Q-Q plot



- A normal Q-Q plot shows if residuals are normally distributed.
- More generally, a Q-Q (quantile-quantile) plot plots two sets of quantiles against one another.
- Quantiles coming from the same distribution should roughly form points along a straight line

Normal Q-Q plot for normal data





Figure: Normal data will tend to produce points on the dashed line of a Normal Q-Q plot.

Normal Q-Q plot



- While visual checks of this sort are subjective, it allows us to see at-a-glance if our assumption is plausible, and if not, how the assumption is violated and what data points contribute to the violation.
- This is not a formal test unlike other **statistical approaches**
- If points fall very far from a straight line, that is cause for concern.

Normal Q-Q plot for skewed data





Figure: Curved Normal Q-Q plots may be an indication that your data are skewed.

Normal Q-Q plot for heavy-tailed data





Figure: Normal Q-Q plots that are flat in the middle and highly curved at extremities may indicate heavy-tailed data.

Conclusion



- While many of the steps towards fitting a linear regression model are algorithmic, model building is more an art than a science.
- These diagnostics tools are meant to guide you through making a decision, but the decision is ultimately yours.