

The University of British Columbia

Data Science 570 Predictive Modelling

Lab Assignment 3

The TA will demonstrate exercises 1 through 3. You are expected to submit answers to exercise 7. Instructions: Please use a png, pdf or html file for submission.

1. Life Auto Data

Today we'll be learning how to plot and interpret several diagnostic plots. To begin, let's start by fitting a linear regression model to the Auto dataset available in the ISLR package. For more details see ?Auto. fit1 attempts to explain gas mileage (in mpg) using the predictor variable of the engine's horsepower.

```
install.packages('ISLR', repos='http://cran.us.r-project.org')

## Installing package into 'C:/Users/xshi/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)

## package 'ISLR' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\xshi\AppData\Local\Temp\RtmpW6fYm4\downloaded_packages

library(ISLR)

## Warning: package 'ISLR' was built under R version 4.0.5

attach(Auto)
head(Auto)

##      mpg cylinders displacement horsepower weight acceleration year origin
## 1    18         8          307         130   3504          12.0     70      1
## 2    15         8          350         165   3693          11.5     70      1
## 3    18         8          318         150   3436          11.0     70      1
## 4    16         8          304         150   3433          12.0     70      1
## 5    17         8          302         140   3449          10.5     70      1
## 6    15         8          429         198   4341          10.0     70      1
##
##              name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5      ford torino
## 6    ford galaxie 500

fit1 <- lm(mpg~horsepower)
summary(fit1)

##
## Call:
```

```
## lm(formula = mpg ~ horsepower)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.5710  -3.2592  -0.3435   2.7630  16.9240
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 39.935861   0.717499   55.66  <2e-16 ***
## horsepower  -0.157845   0.006446  -24.49  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.906 on 390 degrees of freedom
## Multiple R-squared:  0.6059, Adjusted R-squared:  0.6049
## F-statistic: 599.7 on 1 and 390 DF,  p-value: < 2.2e-16
```

Judging by the very low p-values for β_1 (slope), it appears that there is a relationship between the predictor and the response. Since the sign for slope is negative ($\hat{\beta}_1 = -0.1578$) we expect this relationship to be negative, that is, as the horsepower of our engine increases, gas mileage tends to go down. This relationship would also appear to be relatively strong as indicated by the high R^2 value = 0.6049 (ie approximately 60% of the variance in the response is being explained by the simple model).


2. Residual analysis

While the conclusions we made above may seem appropriate based on the R output, we would be amiss to not check the assumptions of this model before we jump to any conclusions. Recall that the assumptions for a linear regression model are that:

- There exists an (at least approximate) linear relationship between Y and X.
- The distribution of ε_i has constant variance.
- ε_i is normally distributed.
- ε_i are independent of one another. For example, ε_2 is not affected by ε_1 .

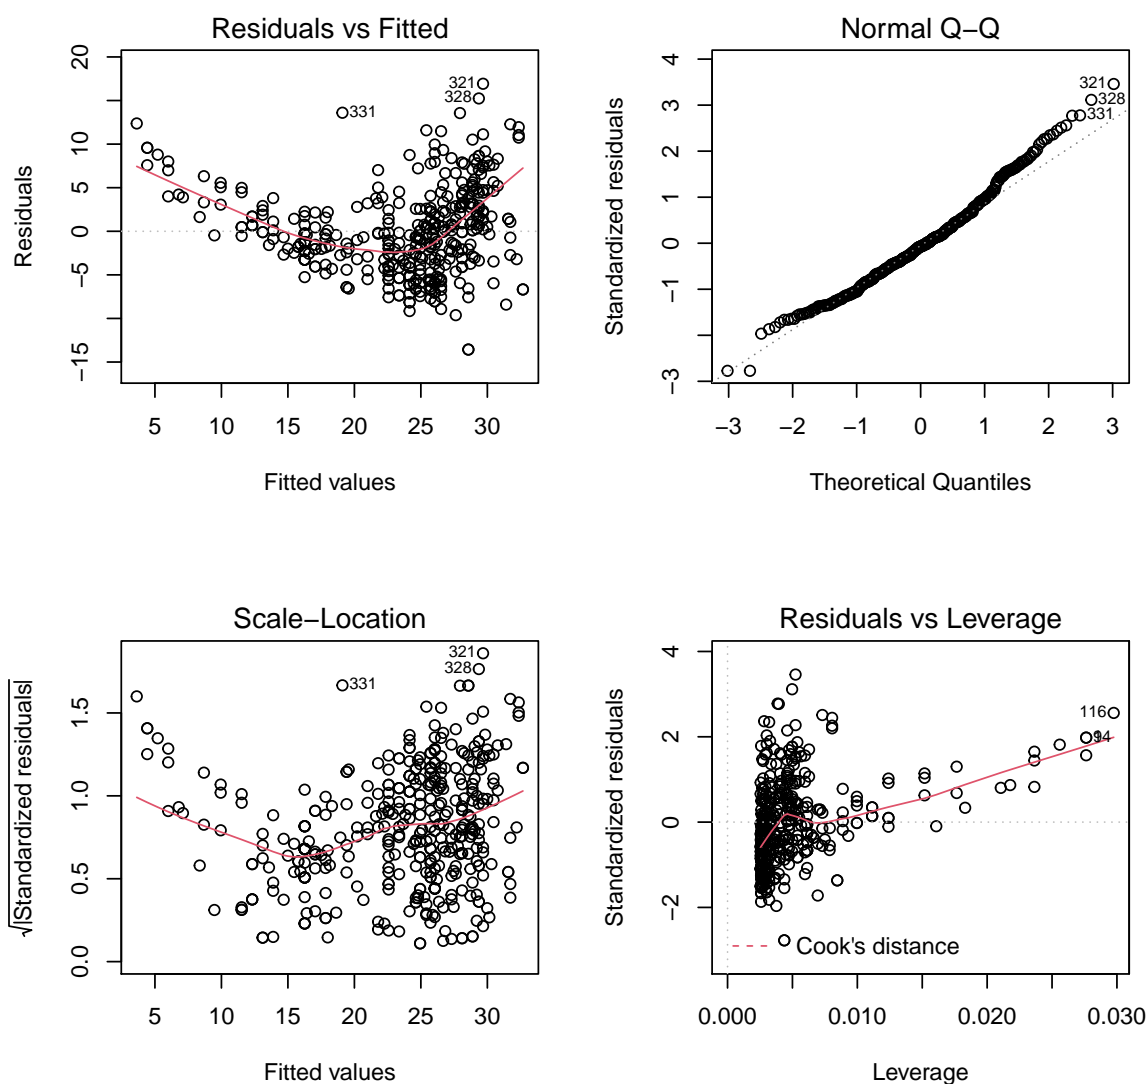
To check the first assumption, we may like to look at the scatterplot with the line of best fit. As demonstrated below, it would appear that the response and predictor variable have a curved relationship rather than a linear one.

```
plot(mpg~horsepower, data=Auto)
abline(fit1, col="red")
```



We could also explore the diagnostic plots produced when call `plot()` on the output of an `lm()` object. To see them all in one plot, let's change the panel layout of our plot window to store 4 figures (in this case in a 2 by 2 matrix) (for more details see `?par`). In other words we could call:

```
par(mfrow=c(2,2))  
plot(fit1)
```



As seen in the Scale-Location graph, it appears there may also been some outliers in our data. Using the ISLR recommendation, we can check the observations whose studentized residuals and observe which ones are greater than 3.

```
which(rstudent(fit1)>3)

## 321 328
## 321 328

## 321 328
```

It appears from the Residuals vs Leverage graph that a number of points could be high leverage points. To investigate the top, say 5, observations having the highest cooks distance, we could do the following:

```
order(cooks.distance(fit1), decreasing = TRUE)[1:5]
```

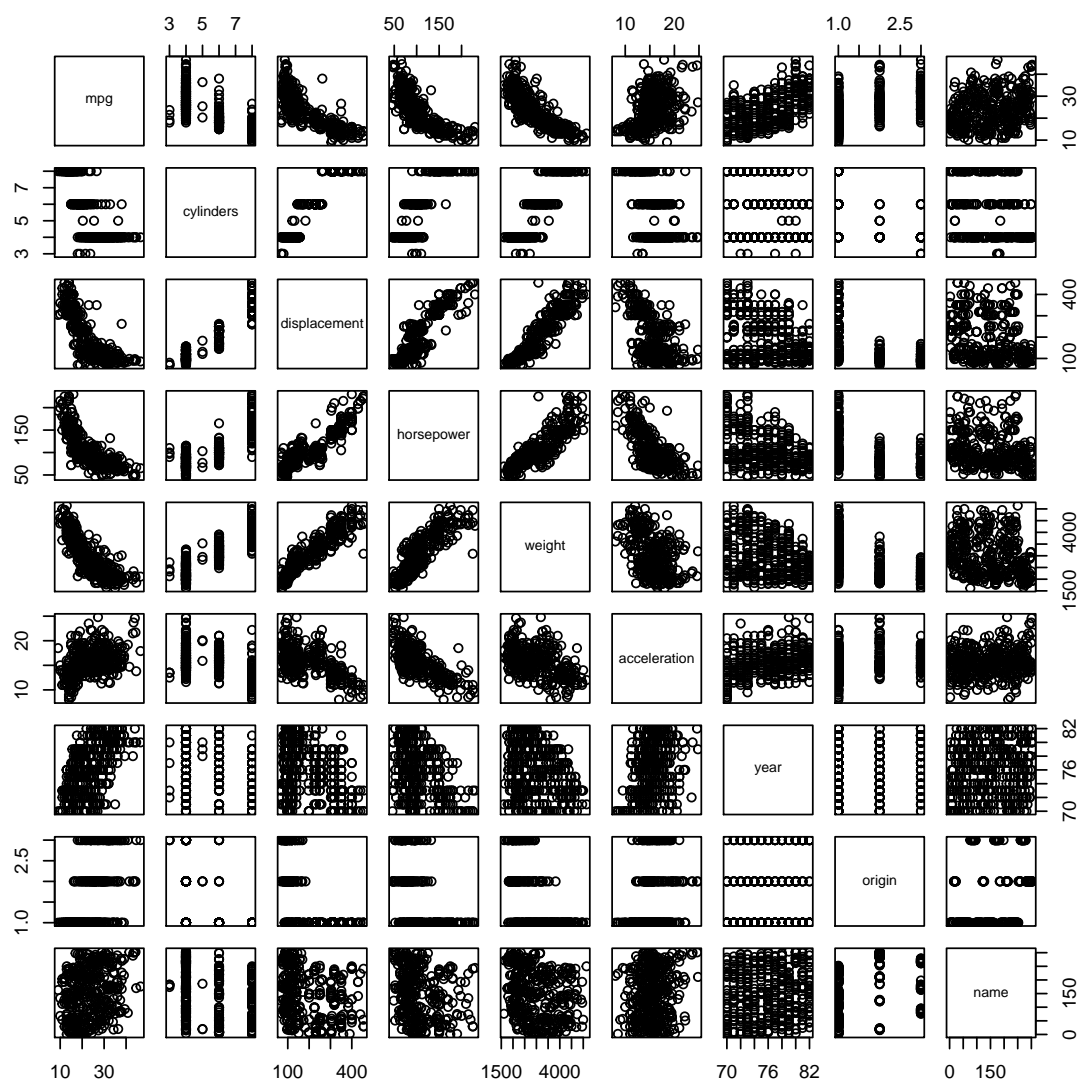
```
## [1] 116  9 14  7 95
```

```
## 116  9 14  7 95
```

3. Collinearity

Now let's consider a MLR model using more predictors from the Auto data from above. We discussed in lecture how some of these predictor variables might be correlated with one another. To investigate this, let's have a look at a scatterplot matrix which includes all of the variables in the data set using the `pairs()` function.

```
pairs(Auto)
```



We can see that horsepower and displacement for example are highly correlated. We might also decide to look at the matrix of correlations between the variables using the function `cor()`.

```
# remove "name" (since it is not numeric)
cor(Auto[,names(Auto)!="name"])

##           mpg  cylinders displacement horsepower      weight
## mpg      1.0000000 -0.7776175   -0.8051269 -0.7784268 -0.8322442
## cylinders -0.7776175  1.0000000    0.9508233  0.8429834  0.8975273
## displacement -0.8051269  0.9508233    1.0000000  0.8972570  0.9329944
## horsepower -0.7784268  0.8429834    0.8972570  1.0000000  0.8645377
## weight     -0.8322442  0.8975273    0.9329944  0.8645377  1.0000000
## acceleration 0.4233285 -0.5046834   -0.5438005 -0.6891955 -0.4168392
## year        0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
## origin      0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
##           acceleration      year      origin
## mpg      0.4233285  0.5805410  0.5652088
## cylinders -0.5046834 -0.3456474 -0.5689316
## displacement -0.5438005 -0.3698552 -0.6145351
## horsepower -0.6891955 -0.4163615 -0.4551715
## weight     -0.4168392 -0.3091199 -0.5850054
## acceleration 1.0000000  0.2903161  0.2127458
## year        0.2903161  1.0000000  0.1815277
## origin      0.2127458  0.1815277  1.0000000
```

Notice that the correlation between horsepower and displacement is very high (0.8972570) as is the correlation between cylinders and displacement is very high (0.9508233). As discussed in lecture, we can go a step further than the bivariate correlations and compute the so-called VIF values. The car package will help us do this.

```
require(car)

## Loading required package: car
## Loading required package: carData

## Loading required package: car
## Loading required package: carData
lm.fit <- lm(mpg ~ .-name, data=Auto)
summary(lm.fit)

##
## Call:
## lm(formula = mpg ~ . - name, data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.5903 -2.1565 -0.1169  1.8690 13.0604
##
```

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -17.218435   4.644294  -3.707  0.00024 ***
## cylinders   -0.493376   0.323282  -1.526  0.12780
## displacement  0.019896   0.007515   2.647  0.00844 **
## horsepower   -0.016951   0.013787  -1.230  0.21963
## weight       -0.006474   0.000652  -9.929 < 2e-16 ***
## acceleration  0.080576   0.098845   0.815  0.41548
## year         0.750773   0.050973  14.729 < 2e-16 ***
## origin       1.426141   0.278136   5.127 4.67e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.328 on 384 degrees of freedom
## Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
## F-statistic: 252.4 on 7 and 384 DF,  p-value: < 2.2e-16

vif(lm.fit)

##      cylinders displacement    horsepower      weight acceleration      year
## 10.737535    21.836792      9.943693    10.831260      2.625806    1.244952
##      origin
## 1.772386
```

4. Consider the simple regression model fit to the lawn roller data.

```
install.packages('DAAG', repos='http://cran.us.r-project.org')

## Installing package into 'C:/Users/xshi/Documents/R/win-library/4.0'
## (as 'lib' is unspecified)
## also installing the dependencies 'rbibutils', 'Rdpack'

##
## There are binary versions available but the source versions are later:
##      binary source needs_compilation
## rbibutils 2.2.8 2.2.16                TRUE
## Rdpack    2.3   2.5                  FALSE
## DAAG      1.24 1.25.4                 FALSE
##
## Binaries will be installed
## package 'rbibutils' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\xshi\AppData\Local\Temp\RtmpW6fYm4\downloaded_packages

## installing the source packages 'Rdpack', 'DAAG'

library(DAAG)
```

```
##
## Attaching package: 'DAAG'
## The following object is masked from 'package:car':
##
##      vif

x=roller$weight
y=roller$depression
```

4.1 2 marks Plot the data points and fitted line.

4.2 2 marks Call plot() on the output of an lm() object.

5. (2 marks) Consider the blood pressure data again. Call pairs() and cor() to show the bivariate correlations and compute the so-called VIF vales.
6. (3 marks) Simulate 100 Poisson random numbers using rpois(100, λ), where $\lambda = e^{\beta_0 + \beta_1 x}$, $\beta_0 = 1$, $\beta_1 = -1$, and x is from a standard normal distribution. Use glm() function to fit Poisson regression and print the coefficients.
7. (3 marks) Simulate 100 Bernoulli random numbers using rbinom(100, 1, prob = p), where $p = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$, $\beta_0 = 1$, $\beta_1 = -1$, and x is from a standard normal distribution. Use glm() function to fit logistic regression and print the coefficients.