

Classification and Regression Trees

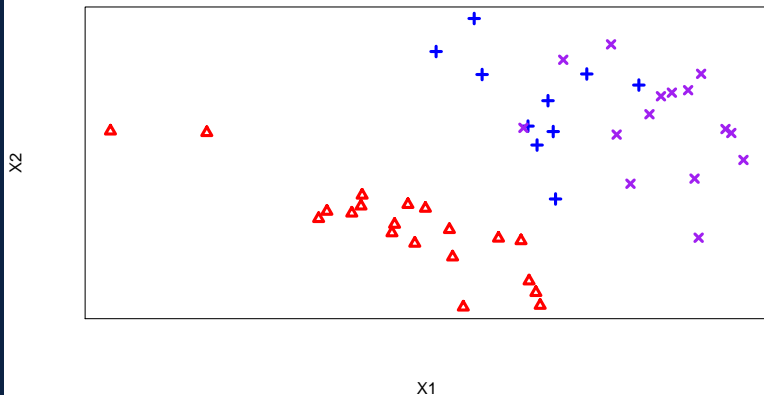
UBCO MDS — DATA 571



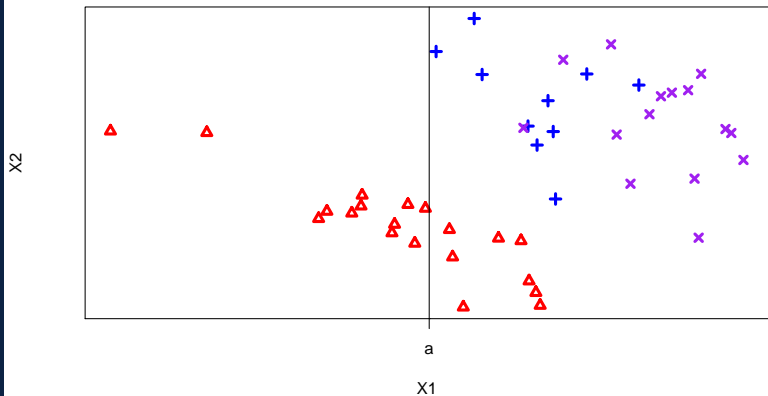


- ▶ Decision trees (or CARTs) partition the predictor space by simple binary splitting rules. Think ‘flow chart’.
- ▶ They have both advantages and disadvantages over other methods...but first, let’s motivate their idea.

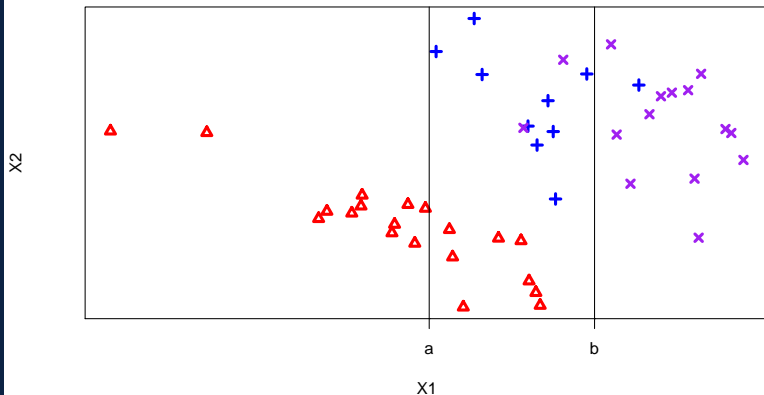
Motivating Example: Classification



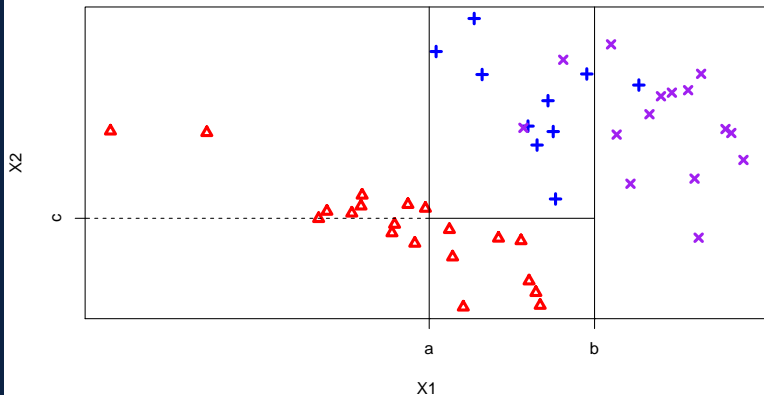
Motivating Example: Classification



Motivating Example: Classification



Motivating Example: Classification




Motivating Example: Classification



- ▶ Are we done with 3 splits?
- ▶ If no then maybe...

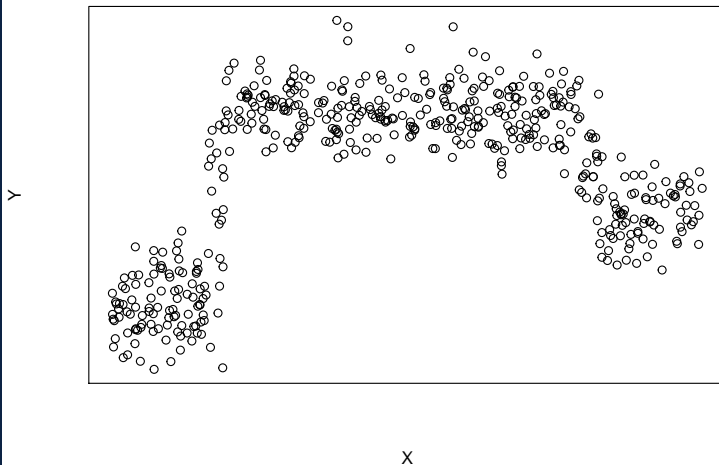
Motivating Example: Classification



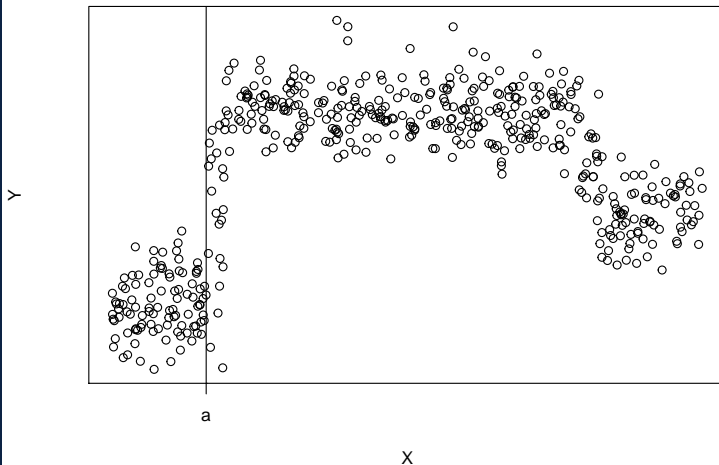


- ▶ So, it appears there are some things we may need to consider going forward in the classification realm.
- ▶ Before diving further into specifics, let's motivate trees for regression through an example.

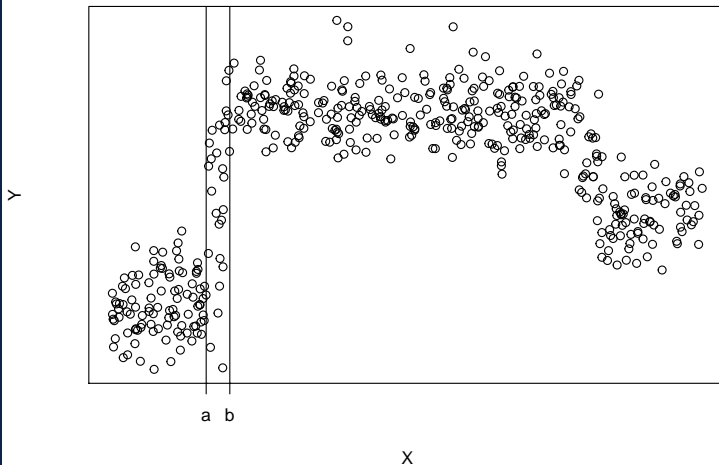
Motivating Example: Regression



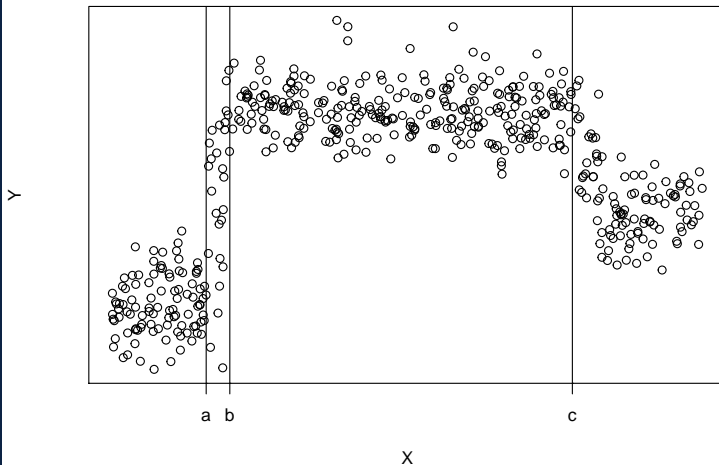
Motivating Example: Regression



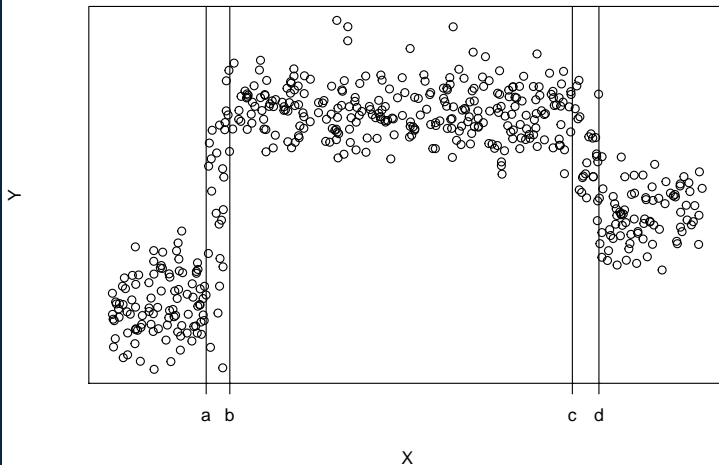
Motivating Example: Regression



Motivating Example: Regression




Motivating Example: Regression





- ▶ Now that we have regions of relative homogeneity in Y ...what would the simplest plan be?
- ▶ Average the observations in the region to provide the predicted value...

Motivating Example: Regression



► Goals:

- Divide the predictor space into J non-overlapping regions R_1, R_2, \dots, R_J . Using binary splits, these regions are hyper-rectangles.
- Regression: Predict the mean response for the observations that fall into each region R_j from the training set.

Classification: Predict the categorical response as the most commonly occurring in each region R_j from the training set.

Growing a Tree

- ▶ How do we 'best' partition the predictor space?
- ▶ For regression: find binary splits which minimize the RSS

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

- ▶ For classification: minimizing the misclassification rate is not usually done at the growth stage. Usually the **Gini index**, **deviance**, or **cross-entropy** is minimized. Most of the time, it won't make a large difference which is chosen. Here's Gini for each node m , where K is the number of response classes...

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

Growing a Tree - Regression



- ▶ We run into a common computational hurdle here! We cannot consider all possible hyper-rectangles in our data-set to find the optimal tree...computationally impossible as dimensionality increases.
- ▶ Instead, we can start seek variable j and cutpoint s such that

$$\sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

- ▶ Then, via the same process, we subdivide region R_1 into R_3 and R_4 . Then we subdivide R_2 into R_5 and R_6 , etc...until some stopping criteria is met (for example, each region contains a minimum of 3 observations).

- ▶ Building a tree through the description just given will tend to heavily overfit the training data.
- ▶ A common strategy is to grow a large tree, and then **prune** it back using **cross-validation** and a **cost-complexity** tuning parameter.

- ▶ We introduce a non-negative tuning parameter α .
- ▶ For every value of α there is a subset (T) of our large tree (T_0) such that

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

is minimized. Where $|T|$ is the number of terminal nodes in the tree.

- ▶ Using k -fold cross validation, we both build trees, and prune them back for possible values of α . Each tree and subtree built in this fashion will produce an estimate of the MSE.
- ▶ Pick the α that minimizes the predicted MSE using the above cross-validation. Then prune the original tree (built on the full data) with the chosen α .

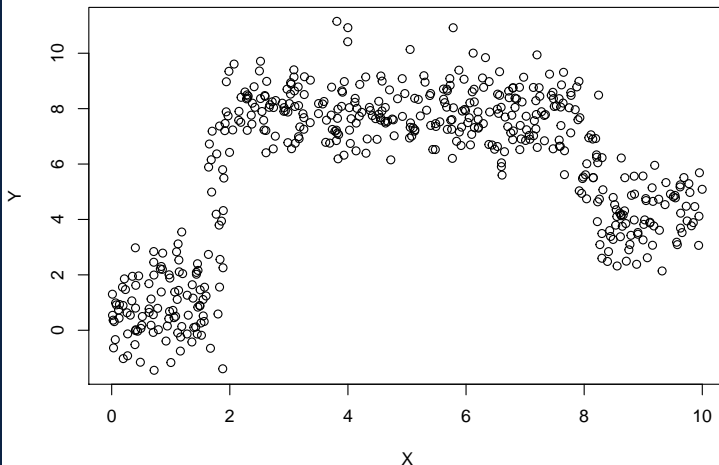


- ▶ For classification trees, the process is very much the same.
- ▶ We use the previously mentioned Gini index for choosing splits (growing the tree), and often use the misclassification rate during the pruning stage (again with cost-complexity tuning parameter).

Motivating Example: Regression



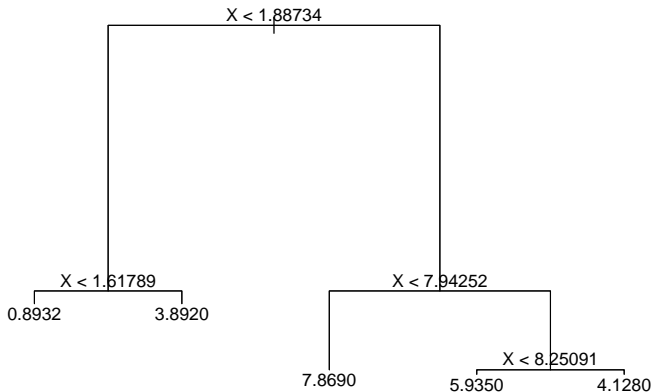
- Back to the motivating example for regression.



Motivating Example: Regression




- Actual fitting in R gives



Motivating Example: Regression




- Which back to the plot gives



Motivating Example: Regression



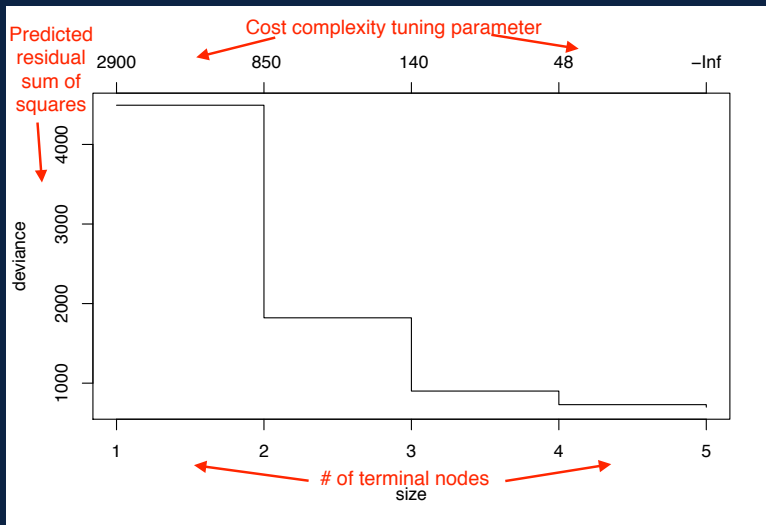
- (Which is very close to what we eyeballed earlier)



Motivating Example: Regression



- BUT! We haven't pruned yet

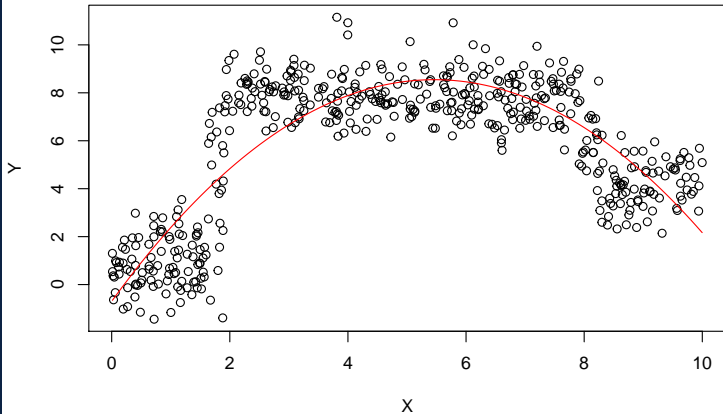


Motivating Example: Regression



- ▶ So it turns out, no pruning needed, according to CV
- ▶ Now, how does it compare to (quadratic) linear regression ?

Motivating Example: Regression



Motivating Example: Regression



- ▶ Training sum squared errors:

- ▶ Tree: 590
- ▶ Quad: 1321

- ▶ CV sum squared errors:

- ▶ Tree: 699
- ▶ Quad: 1345

Example: Regression



- ▶ Let's work with that beer data set again...

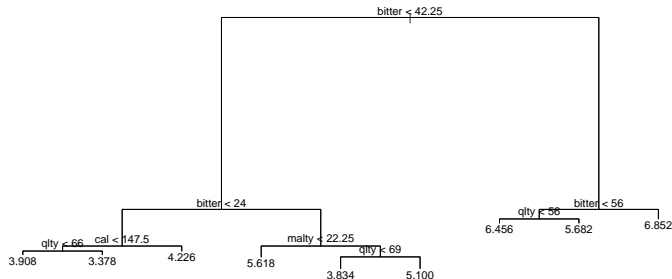
```
> head(beer[,-1])  
  price qlty cal alc bitter malty  
1  6.24  97 159 5.2   52.5  50.5  
2  4.79  92 160 5.0   41.5  48.5  
3  5.96  90 160 4.9   54.5  48.5  
4  4.70  79 162 4.9   37.0  48.0  
5  4.11  72 157 5.5   35.5  41.0  
6  3.85  66 151 4.9   35.0  24.0
```

- ▶ Using a standard call, we fit a large tree...

Example: Regression



```
> library(tree)
> beert <- tree(price~., data=beerdat)
> plot(beert)
> text(beert)
```

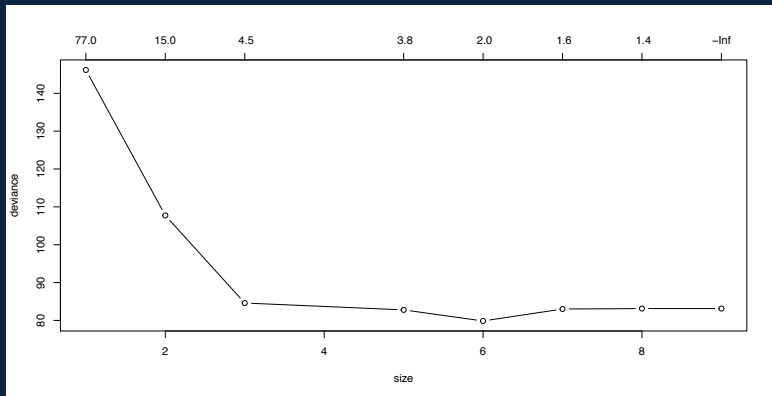


Example: Regression




- ▶ As previously noted, this is likely heavily overfitting, so...

```
> cv.beert <- cv.tree(beert)
> plot(cv.beert, type="b")
```



Example: Regression

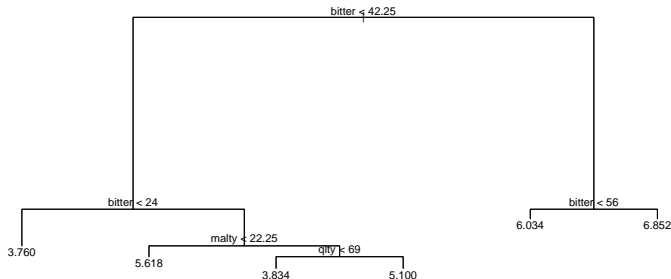


Example: Regression



► Time to prune...

```
> p.beert <- prune.tree(beert, best=6)
> plot(p.beert)
> text(p.beert)
```



Example: Regression



- ▶ We can check the RSS, or equivalently (in this case) the deviance, of the model, and compare to the linear model

```
> sum((yhat-beerdat[,1])^2)
[1] 37.22908
> deviance(p.beert)
[1] 37.22908
> beerlm <- lm(price~., data=beerdat)
> sum(residuals(beerlm)^2)
[1] 47.26182
```

- ▶ So is the regression tree better?

- ▶ Those were training RSS!


```
> library(DAAG)
> beercv <- cv.lm(data=beerdat, beerlm, m=10)
> sum((beercv$price-beercv$cvpred)^2) #10CV RSS for lm
[1] 54.6
> min(cv.beert$dev) #10CV RSS for 6 terminal tree
[1] 74.3
```

- ▶ So even with relatively simple models we need to be careful!

Motivating Example: Classification




- Back to the motivating example for classification.



Motivating Example: Classification



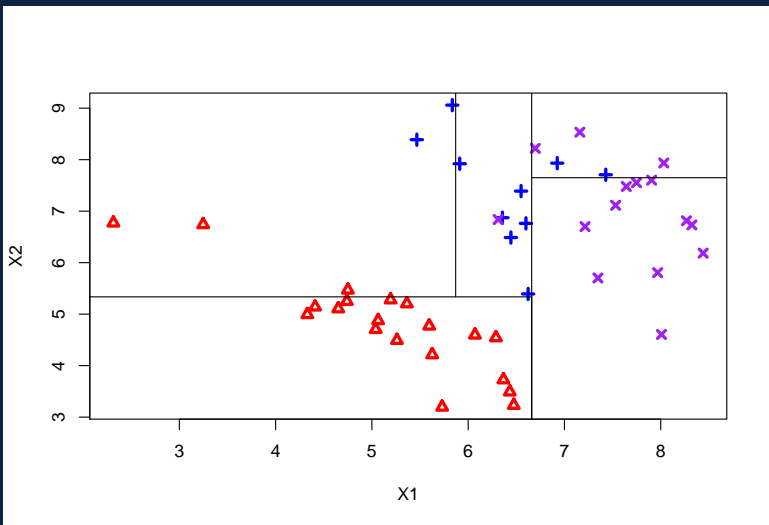
- Under default settings, fitting gives



Motivating Example: Classification



- Results in the following plot. Concerns?



Motivating Example: Classification




- ▶ Default model fitting settings have a minimum of 5 observations in any terminal node.
- ▶ Change the minimum to 3....and we get

Motivating Example: Classification




- Under custom settings, fitting gives



Motivating Example: Classification




- Results in the following plot.



Motivating Example: Classification



- ▶ CV suggests to prune to 4 terminal nodes (for parsimony)



Motivating Example: Classification



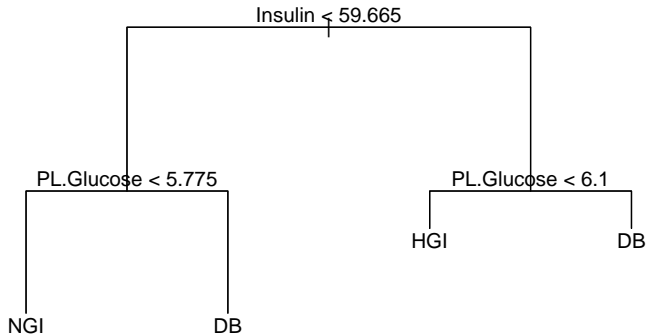
- ▶ Misclassification rate: $3/45 = .067$
- ▶ CV misclassification rate: $12/45 = .267$

Example: Classification




- ▶ Actual prior consultation I worked on.
- ▶ Lots of measurements collected on patients with normal glucose levels (NGI), high levels (HGI), and diabetes (DB)
- ▶ I fit a tree to predict the above classifications, and...

Example: Classification



Example: Classification



Example: Classification



- ▶ Only 2 misclassifications! Rate $2/39 = 0.0513$
- ▶ And, I was told, results match fairly closely to current diagnostic procedures.
- ▶ Cross validation suggests the full size tree proposed in the previous pictures, though it does predict a higher misclassification rate of $10/39 = 0.256$



- ▶ Pros/cons/discussion on board



THE UNIVERSITY OF BRITISH COLUMBIA

