

#### Motivation



- ► We've now seen a couple of alternatives to the linear model for regression.
- ▶ BUT, the linear model still reigns supreme in most realms of science and industry due to its simplicity (which helps for inference).
- Let's see how we can improve upon the linear model.

#### Variable Selection



- ► You have seen (570) forward, backward, and mixed stepwise selection for getting rid of useless predictors.
- ➤ Those methods were largely based on the *p*-values resulting from the *t*-tests on the coefficients...certainly not ideal, especially for large predictor sets.
- ▶ The model measures discussed in 570 (adjusted R², AIC, BIC, etc) can help choose between models, but not suggest which models to compare.
- ► Let's start looking at some more modern methods to intertwine these goals...

Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 3 / 41

# Ridge Regression



► Recall: Least squares estimation minimizes the RSS

$$RSS = \sum_{i=1}^{n} (y_i - \hat{y}_i)^2 = \sum_{i=1}^{n} (y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij})^2$$

► Ridge regression includes a penalty in the estimation process, instead minimizing

$$\sum_{i=1}^{n} (y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij})^2 + \lambda \sum_{j=1}^{p} \beta_j^2 = RSS + \lambda \sum_{j=1}^{p} \beta_j^2$$

where  $\lambda$  is a tuning parameter.

▶ This penalty shrinks the  $\beta_j$  estimates towards 0.





- ▶ If  $\lambda = 0$  then the estimation process is simply least squares once again.
- As  $\lambda \to \infty$  then the penalty grows and the coefficients approach (but never equal) 0.
- ightharpoonup Clearly, some care is required in determining  $\lambda$ . Before diving into that further, let's introduce another (more popular) method of penalizing least squares...

#### Lasso



- ► The 'least absolute shrinkage and selection operator' or lasso is arguably the most popular modern method applied to linear models. It is quite similar to ridge regression...
- ▶ The lasso minimizes

$$\sum_{i=1}^{n} (y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij})^2 + \lambda \sum_{j=1}^{p} |\beta_j| = RSS + \lambda \sum_{j=1}^{p} |\beta_j|$$

where  $\lambda$  is a tuning parameter.

► It turns out that this penalty has an important advantage: it will force some coefficients to 0, whereas RR will only force them 'close' to 0.

# $\lambda$ again



- lacktriangle Both the lasso and ridge regression require specification of  $\lambda$
- ▶ In fact, for each value of  $\lambda$  we will see different coefficients. So how do we find the best one?
- ► The answer, as usual, is cross-validation.



- This approach (especially lasso) is popular with large data sets to simultaneously model in a linear fashion and perform variable selection.
- ► BUT, even if variable selection is not your goal, shrinkage methods are still useful!
- Let's see why...

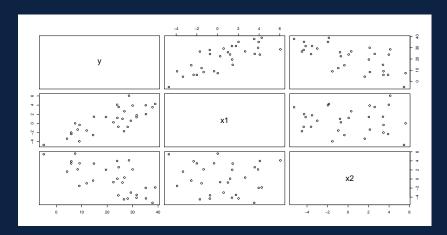


We generate n = 30 data under the following model

$$Y = 20 + 3X_1 - 2X_2 + \epsilon$$

Both  $X_1$  and  $X_2$  are sampled independently as normal.  $\epsilon$  is sampled as N(0, 4).







```
> summary(lm(y~x1+x2))
Call:
lm(formula = y x1 + x2)
Residuals:
   Min 10 Median 30
                               Max
-4.7524 -0.9086 -0.1260 1.7005 3.0130
Coefficients:
          Estimate Std. Error t value Pr(>|t|)
(Intercept) 19.4425 0.3636 53.47 <2e-16 ***
x1
         3.0808 0.1335 23.07 <2e-16 ***
     -2.1267 0.1097 -19.39 <2e-16 ***
x2
Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1
```

```
Residual standard error: 1.941 on 27 degrees of freedom Multiple R-squared: 0.9734, Adjusted R-squared: 0.9714 F-statistic: 494.4 on 2 and 27 DF, p-value: < 2.2e-16
```

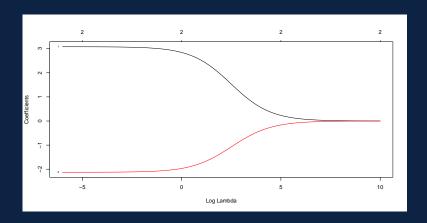


► So, the standard approach gives us the estimated model

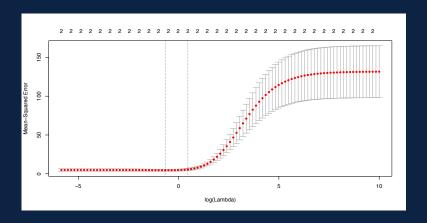
$$\hat{Y} = 19.44 + 3.08X_1 - 2.12X_2$$

- As compared to the true model, we are underestimating the intercept, and overestimating (in absolute value sense) the coefficients
- ➤ So, let's look at ridge regression. Firstly, we need to determine the tuning parameter. Some built-in functions give us the following plots

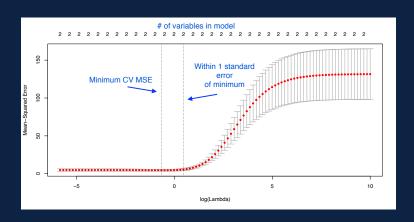














```
> coef(rrsimmin)
3 x 1 sparse Matrix of class "dgCMatrix"
                    s0
(Intercept) 19.460996
x1
             3.045163
x2
            -2.102821
> coef(rrsim1se)
3 x 1 sparse Matrix of class "dgCMatrix"
                    s0
(Intercept) 19.581891
x1
            2.811993
x2
            -1.946284
```



▶

$$\lambda = 0.000$$
  $\hat{Y} = 19.44 + 3.08X_1 - 2.12X_2$   
 $\lambda = 0.141$   $\hat{Y} = 19.46 + 3.04X_1 - 2.10X_2$   
 $\lambda = 1.152$   $\hat{Y} = 19.58 + 2.81X_1 - 1.94X_2$ 

► As compared to the true model,

$$Y = 20 + 3X_1 - 2X_2 + \epsilon$$

the closest coefficient estimations are (arguably) at  $\lambda=0.141$ .

► This is 'proof' that even in simple, low dimensional cases, shrinkage estimators can be useful!

### Example: Beer Data



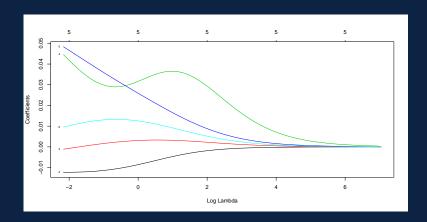
```
> summary(lm(price~., data=beer[,-1]))
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 3.7174521 0.6807382 5.461 8.56e-07 ***
qlty -0.0111435 0.0064446 -1.729 0.0887.
cal
         -0.0055034 0.0089957 -0.612 0.5429
alc 0.0764520 0.1752318 0.436 0.6641
bitter 0.0677117 0.0153219 4.419 3.98e-05 ***
malty 0.0002832 0.0099639 0.028 0.9774
Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1
```

Residual standard error: 0.8661 on 63 degrees of freedom Multiple R-squared: 0.6678, Adjusted R-squared: 0.6415 F-statistic: 25.33 on 5 and 63 DF, p-value: 6.588e-14

RR and LASSO Lecture 6

### Example: Beer Data - RR

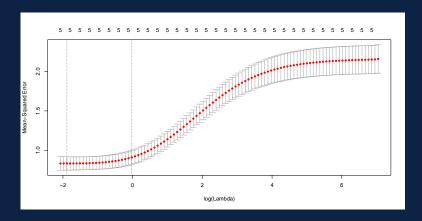




Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 19 / 41

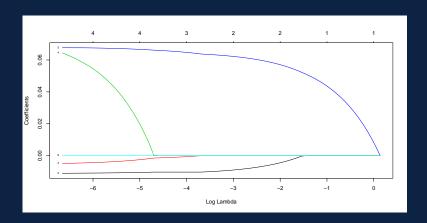
### Example: Beer Data - RR





### Example: Beer Data - lasso

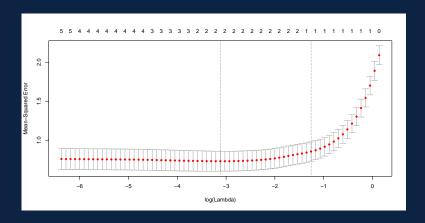




Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 21 / 41

### Example: Beer Data - lasso





Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 22 / 41

## Example: Beer Data Summary



```
LM
                                        R.R.
                                                   LASSO
(Intercept)
              3.7174521392 3.6868302681 3<u>.327790179</u>
qlty
             -0.0111434572 -0.0122442317 -0<u>.008507677</u>
cal
             -0.0055034319 -0.0009067841
alc
              0.0764519623
                             0.0426635115
              0.0677117274 0.0474364378
bitter
                                            0.061693640
malty
              0.0002831604
                             0.0098597370
```

▶ While the true model is unknown, we do know a reduced model is more useful for this data (from previous analyses).

Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 23 / 41

#### **Benefits**

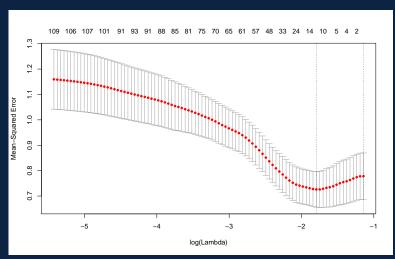


- Examples shown thus far are relatively uncommon examples for RR/LASSO because they do not showcase the core benefits of regularization
- lacktriangle Recall (perhaps) that least squares cannot be fit when p>n
- No longer the case with regularization in effect! So LASSO is a common technique for high-dimensional data sets.
- ► Relatedly, the removal of variables takes care of multicollinearity problems.

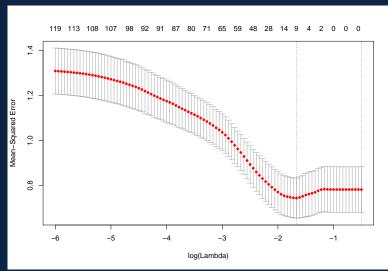


```
> bmat <- matrix(rnorm(50000), nrow=100)
> dim(bmat)
[1] 100 500
> y <- rnorm(100)
> bsimcv <- cv.glmnet(bmat, y, alpha=1)
> plot(bsimcv)
```

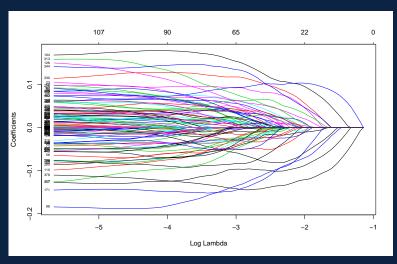












### More Thoughts



- ► Some questions have arisen over the years…let's take a deeper dive on:
- ► Why does RR lead to no variables removed, while LASSO does remove variables?
- ► What is the effect of scaling on the variables?

#### Variable Removal



Firstly, note that the RR optimization, minimizing

$$\sum_{i=1}^{n} (y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij})^2 + \lambda \sum_{j=1}^{p} \beta_j^2$$

could be equivalently written as the standard RSS optimization, minimizing RSS

$$\sum_{i=1}^{n} (y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij})^2$$

subject to  $\sum_{j=1}^{p} \beta_{j}^{2} \leq t$  where t is some positive tuning parameter.

- lacktriangle For LASSO, the constraint would be adjusted to  $\sum_{j=1}^p |eta_j| \leq t$
- ► This formulation will aid visualizations, but may also aid your understanding of regularization in general.

#### Variable Removal

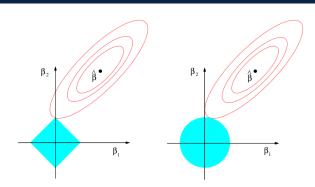


- Now, why does this change in constraint for LASSO lead to variable removal?
- ► Hopefully the interactive visualizations in the app will shed some light on this!!

#### Variable Removal



#### From ESL...



**FIGURE 3.11.** Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions  $|\beta_1| + |\beta_2| \le t$  and  $\beta_1^2 + \beta_2^2 \le t^2$ , respectively, while the red ellipses are the contours of the least squares error function.



- ▶ While it is true that you don't need to worry about scaling when using glmnet in R (it auto-standardizes and retransforms back), it's worth exploring why that automation is important.
- It is relatively easy to see why scaling DOES have an effect on LASSO and RR, and why implementations will generally auto-scale.
- ► The penalty terms, say  $\sum_{j=1}^{p} \beta_j^2 \le t$ , describe a 'fair' relationship between coefficients, assuming the coefficients are on similar scales...
- ► Let's walk through this a bit



```
> set.seed(35521)
> x1 <- rnorm(30)
> x2 <- rnorm(30)
> y <- x1 + x2 + rnorm(length(x1), sd=0.025)</pre>
```

Jeffrey L. Andrews Lecture 6 DATA 571 RR and LASSO 34 / 41



Note that the RSS term is scale invariant on the predictors. > summary(lm(y~x1+x2))

```
Call:
lm(formula = y ~ x1 + x2)
Residuals:
```

```
Min 1Q Median 3Q Max -0.059405 -0.011357 -0.001575 0.021256 0.037663
```

#### Coefficients:

Residual standard error: 0.02426 on 27 degrees of freedom Multiple R-squared: 0.9996, Adjusted R-squared: 0.9996 F-statistic: 3.372e+04 on 2 and 27 DF, p-value: < 2.2e-16



```
> x22 <- x2/10
> summary(lm(y~x1+x22))
```

#### Call:

lm(formula = y ~ x1 + x22)

#### Residuals:

Min 1Q Median 3Q Max -0.059405 -0.011357 -0.001575 0.021256 0.037663

#### Coefficients:

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.02426 on 27 degrees of freedom Multiple R-squared: 0.9996, Adjusted R-squared: 0.9996 F-statistic: 3.372e+04 on 2 and 27 DF, p-value: < 2.2e-16



```
> x12 <- x1/10
> summary(lm(y~x12+x2))
```

#### Call:

```
lm(formula = y ~ x12 + x2)
```

#### Residuals:

```
Min 1Q Median 3Q Max -0.059405 -0.011357 -0.001575 0.021256 0.037663
```

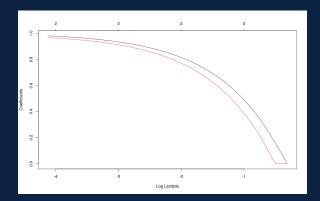
#### Coefficients:

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.02426 on 27 degrees of freedom Multiple R-squared: 0.9996, Adjusted R-squared: 0.9996 F-statistic: 3.372e+04 on 2 and 27 DF, p-value: < 2.2e-16

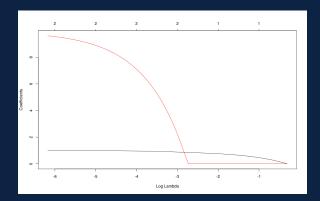


- > run1 <- cv.glmnet(cbind(x1,x2), y, standardize=FALSE)</pre>
- > plot(run1\$glmnet.fit, "lambda")





- > run2 <- cv.glmnet(cbind(x1,x22), y, standardize=FALSE)</pre>
- > plot(run2\$glmnet.fit, "lambda")





- > run3 <- cv.glmnet(cbind(x12,x2), y, standardize=FALSE)</pre>
- > plot(run3\$glmnet.fit, "lambda")

