

Data Manipulation With Pandas Hierarchical Indexing

UBCO Master of Data Science – DATA 542



Motivation

<https://open.canada.ca/data/en/dataset/f140e1d0-ed60-4818-85d6-f02fcb69fda1>

Take a look at the following file:

Wapusk_NP_Coastal_Marine_Snowpack_2006-2016_Data

Station location	Ecozone	Site code	Type of ecosystem	Year	Sample number	Depth (cm)	
Emplacement de la station	Écozone	Code de site	Type d'écosystème	Année	Numéro d'échantillon	Profondeur (cm)	
Mary Lake	Forest	MLK	forest	2006	1	56	
Mary Lake	Forest	MLK	forest	2006	2	38	
Mary Lake	Forest	MLK	forest	2006	3	46	
Mary Lake	Forest	MLK	forest	2006	4	44	
Mary Lake	Forest	MLK	forest	2006	5	41	
Mary Lake	Forest	MLK	forest	2006	6	34	
Mary Lake	Forest	MLK	forest	2006	7	34	
Mary Lake	Forest	MLK	forest	2006	8	38	
Mary Lake	Forest	MLK	forest	2006	9	38	
Mary Lake	Forest	MLK	forest	2006	10	33	
Mary Lake	Forest	MLK	forest	2006	11	45	
Mary Lake	Forest	MLK	forest	2006	12	37	
Mary Lake	Forest	MLK	forest	2006	13	47	
Mary Lake	Forest	MLK	forest	2006	14	42	
Mary Lake	Forest	MLK	forest	2006	15	42	
Mary Lake	Forest	MLK	forest	2006	16	48	
Mary Lake	Forest	MLK	forest	2006	17	56	
Mary Lake	Forest	MLK	forest	2006	18	50	
Mary Lake	Forest	MLK	forest	2006	19	64	

Motivation cont.

Multiple sources of data

Combine or join data

Make distinct analysis levels

Hierarchical indexing

State	Year	population
California	2000	33871648
California	2010	37253956
New York	2000	18976457
New York	2010	19378102
Texas	2000	20851820
Texas	2010	25145561

State	Year	population
California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

Discuss: How will you analyze this data?
e.g. calculate the population in each year, or by each state?

First solution

```
index = [('California', 2000), ('California',  
2010), ('New York', 2000), ('New York', 2010),  
('Texas', 2000), ('Texas', 2010)]
```

```
populations = [33871648, 37253956, 18976457,  
19378102, 20851820, 25145561]
```

```
pop = pd.Series(populations, index=index)
```

```
pop
```

Discuss: 1) pros and cons of this solution
2) Select values for 2010 only

Better solution: Hierarchical indexing

```
index = pd.MultiIndex.from_tuples(index)
pop = pop.reindex(index)
```

Blank entries means the
same value as above



This is an multiply indexed Series

California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

MultiIndex is like an extra dimension

Discuss: Select values for 2010 only

Indexing a MultiIndex Series

Index the individual elements by multiple terms:

```
pop[ 'California', 2000]
```

Partial indexing: Indexing on just one of the levels

The outer level:

```
pop[ 'California' ]
```

The inner level: pass an empty index for the outer level:

```
pop[ :, 2010]
```


Slicing a MultiIndex Series

Slicing can only be done on **sorted indices**. Otherwise it will be a key error.

```
pop.loc['California':'New York']
```

The inner level:

```
pop.loc[:, 2000]
```

Using Boolean masks:

```
pop[pop > 220000000]
```

Fancy indexing

```
pop[['California', 'Texas']]
```

Question

Question 1: which one selects the values for 2010 only, for all states?

A) `pop.loc[:, 2010]`

B) `pop[:, 2010]`

C) `pop.loc[2010]`

D) `pop.loc['California', 2010]`

E) A and B

California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

Stacking and unstacking

We can have one of the indices as a column

```
pop_df = pop.unstack()
```

	2000	2010
California	33871648	37253956
New York	18976457	19378102
Texas	20851820	25145561

Or convert any dataframe to multi-indexing

```
pop_df.stack()
```

California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

Flexibility with hierarchical indexing

Represent data of three or more dimensions in a Series or DataFrame

You may add another column

```
pop_df = pd.DataFrame({'total': pop,
    'under18': [9267089, 9284094, 4687374, 4318033, 5906301,
6879014]})
```

Use functions

```
u18_percent = pop_df['under18']/pop_df['total']
u18_percent.unstack()
```

	2000	2010
California	0.273594	0.249211
New York	0.247010	0.222831
Texas	0.283251	0.273568

Create hierarchical DataFrames

1- Pass list of lists as the index or column when creating a DataFrame

```
df = pd.DataFrame(np.random.rand(4, 2),
index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],
columns=['data1', 'data2'])
```

2- Use tuples as the Keys when using Dictionaries to create a DataFrame

```
data = {('California', 2000): 33871648,
('California', 2010): 37253956, ('Texas', 2000):
20851820, ('Texas', 2010): 25145561, ('New York',
2000): 18976457, ('New York', 2010): 19378102}

pd.Series(data)
```

Create hierarchical DataFrames cont.

3- Explicitly use MultiIndex constructor

```
pd.MultiIndex.from_arrays([['a', 'a', 'b',  
'b'], [1, 2, 1, 2]])
```

```
pd.MultiIndex.from_tuples([('a', 1), ('a', 2),  
('b', 1), ('b', 2)])
```

From Cartesian product of single indices

```
pd.MultiIndex.from_product([['a', 'b'], [1,  
2]])
```

Naming index levels

Name the indexes in hierarchical indexing

```
pop.index.names = ['state', 'year']
```

This is a **name** and not a **label**. So, partial indexing is through the levels attribute of the `MultiIndex` object. For example, it is an error: `pop.loc['state']` while you can run

```
pop.loc['California']
```

state	year	
California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

Try it yourself

Columns can have multiple levels of indexing. Try the following code and discuss slicing or sub-setting the DataFrame

```
# hierarchical indices and columns
index = pd.MultiIndex.from_product([[2013, 2014], [1, 2]],
names=['year', 'visit'])

columns = pd.MultiIndex.from_product(['Bob', 'Guido', 'Sue'],
['HR', 'Temp'], names=['subject', 'type'])

# mock some data
data = np.round(np.random.randn(4, 6), 1)
data[:, :2] *= 10
data += 37

# create the DataFrame
health_data = pd.DataFrame(data, index=index, columns=columns)
health_data
```

Question: Show the results for Bob

Indexing a multiply index DataFrame

The indexing operations on Multiply index Series applies on the columns

```
health_data['Bob']
```

```
health_data['Bob', 'HR']
```

```
health_data.iloc[:2, :2]
```

TRY IT

Each individual index in loc or iloc can be passed a tuple of multiple indices:

```
health_data.loc[:, ('Sue', 'Temp')]
```

```
health_data.loc[(2013, 1), ('Sue', 'Temp')]
```

Do not create slice within a tuple: Syntax errors:

```
health_data.loc[:, 1), (:, 'HR')]
```

Indexing a multiply index DataFrame cont.

Syntax errors: `health_data.loc[:, 1], (:, 'HR')]`

Use python IndexSlice object:

TRY IT

```
idx = pd.IndexSlice
```

```
health_data.loc[idx[:, 1], idx[:, 'HR']]
```

		subject	Bob	Guido	Sue
		type	HR	HR	HR
year	visit				
2013	1	40.0	20.0	39.0	
2014	1	25.0	51.0	49.0	

Sorting hierarchical indexing DataFrame

Sort the data values based on each level of the index labels or column labels using

```
data_frame.sort_index(level = index_level,  
axis= axis)
```



Level starts at 0 indicating the outer index level

Use level to indicate explicitly which level of the indexing to sort

By default, it sorts the indexes.

Use axis to apply sorting on columns.

Sorting hierarchical indexing DataFrame cont.

```
frame = pd.DataFrame(
    np.arange(18).reshape((6, 3)),
    index=[
        ['a', 'a', 'c', 'c', 'b', 'b'],
        [1, 2, 2, 1, 1, 2]],
    columns=[
        ['Ohio', 'Ohio', 'Colorado'],
        ['Green', 'Red', 'Green']])
```

Partial slicing returns key error: `data['a':'b']`

		Ohio		Colorado	
		Green	Red	Green	
a	1	0	1	2	
	2	3	4	5	
c	2	6	7	8	
	1	9	10	11	
b	1	12	13	14	
	2	15	16	17	

Sorting hierarchical indexing DataFrame cont.

```
Frame.sort_index()
```

```
Frame.sort_index(level=0)
```

```
Frame.sort_index(level=1)
```

		Ohio		Colorado	
		Green	Red	Green	
a	1	0	1	2	
	2	3	4	5	
b	1	12	13	14	
	2	15	16	17	
c	1	9	10	11	
	2	6	7	8	

		Ohio		Colorado	
		Green	Red	Green	
a	1	0	1	2	
b	1	12	13	14	
c	1	9	10	11	
a	2	3	4	5	
b	2	15	16	17	
c	2	6	7	8	

Sorting hierarchical indexing DataFrame cont.

```
Frame.sort_index(axis = 1)
```

Use assignment or `inplace = True` to modify the DataFrame not its copy

		Colorado	Ohio	
		Green	Green	Red
a	1	2	0	1
	2	5	3	4
c	2	8	6	7
	1	11	9	10
b	1	14	12	13
	2	17	15	16

Reorder hierarchical indexing DataFrame cont.

```
frame.swaplevel()
frame.swaplevel(0,1)
frame.swaplevel('key1','key2')
```

		Colorado	Ohio	
		Green	Green	Red
key2	key1			
1	a	2	0	1
2	a	5	3	4
1	b	14	12	13
2	b	17	15	16
1	c	11	9	10
2	c	8	6	7

Unstacking multiply index DataFrames

Unstack the hierarchical index DataFrames (reduce dimension) using

`data_frame.unstack(level = level)`

Recover using `stack()`

year	2000	2010
state		
California	33871648	37253956
New York	18976457	19378102
Texas	20851820	25145561

`pop.unstack(level = 1)`

state	year	
California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

state	California	New York	Texas
year			
2000	33871648	18976457	20851820
2010	37253956	19378102	25145561

`pop.unstack(level = 0)`

Rearranging DataFrames

Flattening the hierarchical index DataFrame

Use index as the columns

```
pop.reset_index(name =  
'population')
```

```
flat_pop.set_index(['year',  
'state'])
```

	state	year	population
0	California	2000	33871648
1	California	2010	37253956
2	New York	2000	18976457
3	New York	2010	19378102
4	Texas	2000	20851820
5	Texas	2010	25145561

population		
year	state	
2000	California	33871648
2010	California	37253956
2000	New York	18976457
2010	New York	19378102
2000	Texas	20851820
2010	Texas	25145561

Rearranging DataFrames cont.

Use columns as index

1 `frame.set_index(['c','d'])`

2 `frame.set_index(['c','d'], drop=False)`

	a	b	c	d
0	0	7	one	0
1	1	6	one	1
2	2	5	one	2
3	3	4	two	0
4	4	3	two	1
5	5	2	two	2
6	6	1	two	3

1

	a	b
c d		
one 0	0	7
1 1	6	
2 2	5	
two 0	3	4
1 4	3	
2 5	2	
3 6	1	

2

	a	b	c	d
c d				
one 0	0	7	one	0
1 1	6	one	1	
2 2	5	one	2	
two 0	3	4	two	0
1 4	3	two	1	
2 5	2	two	2	
3 6	1	two	3	

Summary statistics in each level: aggregations

Apply built-in aggregation functions such as `sum()`, `mean()`, `max()` on a particular level and on a specified axis. Axis is required for operations on columns.

```
data_frame.method(level = level_name, axis =  
axis)
```

	subject	Bob		Guido		Sue	
	type	HR	Temp	HR	Temp	HR	Temp
year	visit						
2013	1	114.0	111.0	94.0	111.9	113.0	113.3
	2	118.0	111.2	122.0	110.3	109.0	110.9
2014	1	99.0	111.6	125.0	111.7	123.0	111.4
	2	114.0	111.8	106.0	111.7	91.0	111.3

```
health_data.columns.names
```

```
FrozenList(['subject', 'type'])
```

```
health_data.columns.labels
```

```
FrozenList([[0, 0, 1, 1, 2, 2], [0, 1, 0, 1, 0, 1]])
```

```
health_data.columns.levels
```

```
FrozenList([['Bob', 'Guido', 'Sue'], ['HR', 'Temp']])
```

Summary statistics in each level: aggregations cont.

```
health_data.mean(level='year')
```

subject	Bob		Guido		Sue	
type	HR	Temp	HR	Temp	HR	Temp
year						
2013	116.0	111.1	108.0	111.1	111.0	112.10
2014	106.5	111.7	115.5	111.7	107.0	111.35

```
health_data.sum(axis = 1,
                 level='type')
```

	type	HR	Temp
year	visit		
2013	1	321.0	336.2
	2	349.0	332.4
2014	1	347.0	334.7
	2	311.0	334.8

Learning outcomes

At the end of this lecture you should be able to:

Create MultiIndex objects

Perform indexing and slicing on multiply indexed data

Compute statistics across multiply indexed data

Convert between simple and hierarchically indexed representations of data

Combining and Merging data With Pandas

UBCO Master of Data Science – DATA 542



Combining and merging data

Various sources of data

Merge with:

- 1) `pandas.merge`: connects rows in DataFrames based on one or more keys. Similar to database *join* operations.
- 2) `pandas.concat`: concatenates along an axis.

1) pandas.merge

`pandas.merge()` implements three types of join:

One-to-one

Many-to-one

Many-to-many

`pandas.merge(df_1, df_2, OPTIONS)`

OPTIONS:

- **on:** Column names to join on. Must be found in both dataframes.
- **right_on:** Columns in right DataFrame to use as join keys.
- **left_on:** Columns in left DataFrame to use as join keys.
- **left_index:** Use row index in left as its join key (or keys, if a MultiIndex).
- **right_index:** Use row index in right as its join key (or keys, if a MultiIndex).
- **how:** One of 'inner', 'outer', 'left', or 'right'; defaults to 'inner'.
- **Copy:** If False, avoid copying data into resulting data structure in some exceptional cases; by default always copies.

Example: one-to-one join

Census dataset – BC - 2016

	census_year	geo_name	dissemination_profiles	sex_male
0	2016	British Columbia / Colombie-Britannique	Total - Age groups and average age of the popu...	2278245
1	2016	British Columbia / Colombie-Britannique	0 to 4 years	113355
2	2016	British Columbia / Colombie-Britannique	5 to 9 years	122070
3	2016	British Columbia / Colombie-Britannique	10 to 14 years	119975

	census_year	geo_name	dissemination_profiles	sex_female
0	2016	British Columbia / Colombie-Britannique	Total - Age groups and average age of the popu...	2369810
1	2016	British Columbia / Colombie-Britannique	0 to 4 years	107270
2	2016	British Columbia / Colombie-Britannique	5 to 9 years	114830
3	2016	British Columbia / Colombie-Britannique	10 to 14 years	113885

Merge

```
pandas.merge(df_male, df_female)
```

	census_year	geo_name	dissemination_profiles	sex_male	sex_female
0	2016	British Columbia / Colombie-Britannique	Total - Age groups and average age of the popu...	2278245	2369810
1	2016	British Columbia / Colombie-Britannique	0 to 4 years	113355	107270
2	2016	British Columbia / Colombie-Britannique	5 to 9 years	122070	114830
3	2016	British Columbia / Colombie-Britannique	10 to 14 years	119975	113885
4	2016	British Columbia / Colombie-Britannique	15 to 19 years	133000	125980
5	2016	British Columbia / Colombie-Britannique	20 to 24 years	147615	139945

Joins on the shared columns

Merge cont.

```
pandas.merge(df_male, df_female, on =
['census_year', 'dissemination_profiles'])
```

	census_year	geo_name_x	dissemination_profiles	sex_male	geo_name_y	sex_female
0	2016	British Columbia / Colombie-Britannique	Total - Age groups and average age of the popu...	2278245	British Columbia / Colombie-Britannique	2369810
1	2016	British Columbia / Colombie-Britannique	0 to 4 years	113355	British Columbia / Colombie-Britannique	107270
2	2016	British Columbia / Colombie-Britannique	5 to 9 years	122070	British Columbia / Colombie-Britannique	114830
3	2016	British Columbia / Colombie-Britannique	10 to 14 years	119975	British Columbia / Colombie-Britannique	113885

Joins on the specified keys. If there are shared keys, it prefixes with `_x` and `_y`.

Modify names using `suffixes=(required_name)` attribute:
`suffixes=('_m', '_f')`

Inner join

Merging on the intersection of two DataFrames

```
pd.merge(df6, df7, how='inner')
```

df6

	name	food
0	Peter	fish
1	Paul	beans
2	Mary	bread

df7

	name	drink
0	Mary	wine
1	Joseph	beer

pd.merge(df6, df7)

	name	food	drink
0	Mary	bread	wine

Outer join

An *outer join* returns a join over the union of the input columns, and fills in all missing values with NaNs

```
pd.merge(df6, df7, how='outer')
```

Discuss the results

Left/right join

An *left join* returns a join preserving the entries in the left dataset. It fills the missing values with NaNs. The right join is similar, keeping the entries from the right DataFrame.

```
pd.merge(df6, df7, how='left')
```

Discuss the results

Joining on hierarchical DataFrames

Use multiple keys as the join keys

```
pd.merge(left_df, right_df, left_on=['key1',
'key2'], right_index=True)
```

left_df

	key1	key2	data
0	Ohio	2000	0.0
1	Ohio	2001	1.0
2	Ohio	2002	2.0
3	Nevada	2001	3.0
4	Nevada	2002	4.0

right_df

		event1	event2
Nevada	2001	0	1
	2000	2	3
Ohio	2000	4	5
	2000	6	7
	2001	8	9
	2002	10	11

It is also possible to use indices of both sides

Discuss the results

Question

Question 1: What is the result of the following code?

```
pd.merge(left_df, right_df, left_on=['key1',
'key2'], right_index=True, how='outer')
```

left_df:

	key1	key2	data
0	Ohio	2000	0.0
1	Ohio	2001	1.0
2	Ohio	2002	2.0
3	Nevada	2001	3.0
4	Nevada	2002	4.0

right_df:

		event1	event2
Nevada	2001	0	1
	2000	2	3
Ohio	2000	4	5
	2000	6	7
	2001	8	9
	2002	10	11

A)

	key1	key2	data	event1	event2
0	Ohio	2000	0.0	4.0	5.0
0	Ohio	2000	0.0	6.0	7.0
1	Ohio	2001	1.0	8.0	9.0
2	Ohio	2002	2.0	10.0	11.0
3	Nevada	2001	3.0	0.0	1.0
4	Nevada	2002	4.0	NaN	NaN
4	Nevada	2000	NaN	2.0	3.0

B)

	key1	key2	data	event1	event2
0	Ohio	2000	0.0	4	5
0	Ohio	2000	0.0	6	7
1	Ohio	2001	1.0	8	9
2	Ohio	2002	2.0	10	11
3	Nevada	2001	3.0	0	1

C)

	key1	key2	data	event1	event2
0	Ohio	2000	0.0	4.0	5.0
0	Ohio	2000	0.0	6.0	7.0
1	Ohio	2001	1.0	8.0	9.0
2	Ohio	2002	2.0	10.0	11.0
3	Nevada	2001	3.0	0.0	1.0
4	Nevada	2002	4.0	NaN	NaN

Join attribute

`join` is an attribute of `DataFrame` objects with operations similar to `merge()` to merge on indices.

```
df1.join([df_list], on=key, how=how)
```

- You can pass a list of `DataFrames` or a single `DataFrame`
- `On`: Column names to join on. Must be found in both dataframes.
- `how`: One of 'inner', 'outer', 'left', or 'right'; defaults to 'inner'.

join defaults to left join

Concatenation along axes

Think about:

- 1) “If the objects are indexed differently on the other axes, should we combine the distinct elements in these axes or use only the shared values (the intersection)?”
- 2) “Do the concatenated chunks of data need to be identifiable in the resulting object?”
- 3) “Does the “concatenation axis” contain data that needs to be preserved? In many cases, the default integer labels in a DataFrame are best discarded during concatenation.”

Concatenation along axes cont.

Use `pd.concat()` with the following attributes:

- `obj`: You can pass a list of Pandas objects
- `axis`: The axis along which the join operates
- `keys`: Values to associate with objects being concatenated, forming a hierarchical index along the concatenation axis.
- `how`: One of 'inner', 'outer', 'left', or 'right'; defaults to 'inner'.
- `verify_integrity`: Check new axis in concatenated object for duplicates and raise exception if so; by default (False) allows duplicates.
- `ignore_index`: Do not preserve indexes along concatenation axis, instead producing a new `range(total_length)` index
- `names`: Names for created hierarchical levels if keys and/or levels passed
- `join_axes`: Specific indexes to use for the other $n-1$ axes instead of performing union/intersection logic

Concatenation example

```
s1 = pd.Series([0, 1], index=['a', 'b'])
s2 = pd.Series([2, 3, 4], index=['c', 'd', 'e'])
s3 = pd.Series([5, 6], index=['f', 'g'])
```

a	0
b	1
c	2
d	3
e	4
f	5
g	6

```
pd.concat([s1, s2, s3],
axis=1, sort=False)
```

```
pd.concat([s1, s2, s3])
```

	0	1	2
a	0.0	NaN	NaN
b	1.0	NaN	NaN
c	NaN	2.0	NaN
d	NaN	3.0	NaN
e	NaN	4.0	NaN
f	NaN	NaN	5.0
g	NaN	NaN	6.0

Question

Question 1: What is the result of the following code?

```
s1 = pd.Series([0, 1], index=['a', 'b'])
s4 = pd.Series([0, 1, 5, 6],
index=['a', 'b', 'f', 'g'])
pd.concat([s1, s4], axis=1, sort=True)
```

A)

	0	1
a	0.0	0
b	1.0	1
f	NaN	5
g	NaN	6

B)

	0	1
a	0	0
b	1	1

C)

a	0
b	1
a	0
b	1
f	5
g	6
dtype: int64	

D)

	0	1
a	0.0	0.0
c	NaN	NaN
b	1.0	1.0
e	NaN	NaN

Concatenation through inner join

Intersect them by passing `join='inner'`

```
s1 = pd.Series([0, 1], index=['a', 'b'])
s4 = pd.Series([0, 1, 5, 6], index=['a', 'b', 'f', 'g'])
pd.concat([s1, s4], axis=1, join='inner')
```

	0	1
a	0	0
b	1	1

Specify the axes to be used on the other axes with

```
join_axes:
pd.concat([s1, s4], axis=1,
          join_axes=[['a', 'c', 'b', 'e']])
```

	0	1
a	0.0	0.0
c	NaN	NaN
b	1.0	1.0
e	NaN	NaN

Concatenation issues

Results are not identifiable

Solution: use hierarchical indexing

```
result = pd.concat([s1, s1, s3], keys=['one', 'two', 'three'])
```

one	a	0
	b	1
two	a	0
	b	1
three	f	5
	g	6

Result.unstack:

	a	b	f	g
one	0.0	1.0	NaN	NaN
two	0.0	1.0	NaN	NaN
three	NaN	NaN	5.0	6.0

Concatenation issues cont.

Along axis =1 the keys become column headers

```
result = pd.concat([s1, s1, s3], keys=['one', 'two', 'three'], axis = 1)
```

	one	two	three
a	0.0	0.0	NaN
b	1.0	1.0	NaN
f	NaN	NaN	5.0
g	NaN	NaN	6.0

Result.unstack:

one	a	0.0
	b	1.0
	f	NaN
	g	NaN
two	a	0.0
	b	1.0
	f	NaN
	g	NaN
three	a	NaN
	b	NaN
	f	5.0
	g	6.0

Concatenate DataFrames

Same logic applies to DataFrame concatenation

```
pd.concat([df1, df2], axis=1, keys=['level1',  
'level2'], sort =True, names=['upper', 'lower'])
```

df1

	one	two
a	0	1
b	2	3
c	4	5

df2

	three	four
a	5	6
c	7	8

Create hierarchical index
using keys attribute

	level1		level2	
	one	two	three	four
a	0	1	5.0	6.0
b	2	3	NaN	NaN
c	4	5	7.0	8.0

Name the created axis levels:

upper		level1		level2	
lower		one	two	three	four
a		0	1	5.0	6.0
b		2	3	NaN	NaN
c		4	5	7.0	8.0

Important NOTE

Pandas concatenation *preserves indices*, even if the result will have duplicate indices! Use `verify_integrity` attribute to catch errors

try:

```
pd.concat([df1, df2],
          join_axes=[df1.columns],
          verify_integrity=True, sort = True)
except ValueError as e:
    print("ValueError:", e)
```

df1

df2

	one	two
a	0.0	1.0
b	2.0	3.0
c	4.0	5.0
a	NaN	NaN
c	NaN	NaN

ValueError: Indexes have overlapping values: Index(['a', 'c'], dtype='object')


Important NOTE

If indices are not important, turn the flag `ignore_index` on to create new indices for the concatenated DataFrame

```
pd.concat([df1, df2], join_axes=[df1.columns],
ignore_index=True)
```

	one	two
a	0.0	1.0
b	2.0	3.0
c	4.0	5.0

a	NaN	NaN
c	NaN	NaN



	one	two
0	0.0	1.0
1	2.0	3.0
2	4.0	5.0
3	NaN	NaN
4	NaN	NaN

Alternation and optimization

Use `append` instead of `concat()` : `df1.append(df2)`

Append method in Pandas

- Does not modify the original object
- Is not an efficient method

Optimization:

Build a list of DataFrames

Pass them all at once to the `concat()` function.

Combining data with overlap

You can combine the NaN values of a DataFrame with values from another DataFrame

```
df1.combine_first(df2)
```

```
df1.combine_first(df2)
```

	a	b	c
0	1.0	NaN	2
1	NaN	2.0	6
2	5.0	NaN	10
3	NaN	6.0	14

df1

	a	b
0	5.0	NaN
1	4.0	3.0
2	NaN	4.0
3	3.0	6.0
4	7.0	8.0

df2

	a	b	c
0	1.0	NaN	2.0
1	4.0	2.0	6.0
2	5.0	4.0	10.0
3	3.0	6.0	14.0
4	7.0	8.0	NaN

Similar operation in numpy is where:
`np.where(pd.isnull(a), b, a)`
 It uses b values for a entries that are null

Data Grouping and Aggregation

UBCO Master of Data Science – DATA 542



Introduction

Gain more insights to your data

Summarize data

Statistics of data

Pandas objects aggregation method

Aggregation methods for Pandas DataFrame and Series objects:

Aggregation	Description
count()	Total number of items
first(), last()	First and last item
mean(), median()	Mean and median
min(), max()	Minimum and maximum
std(), var()	Standard deviation and variance
mad()	Mean absolute deviation
prod()	Product of all items
sum()	Sum of all items

Aggregations

Aggregation functions:

`sum()`, `min()`, `max()`, `mean()`

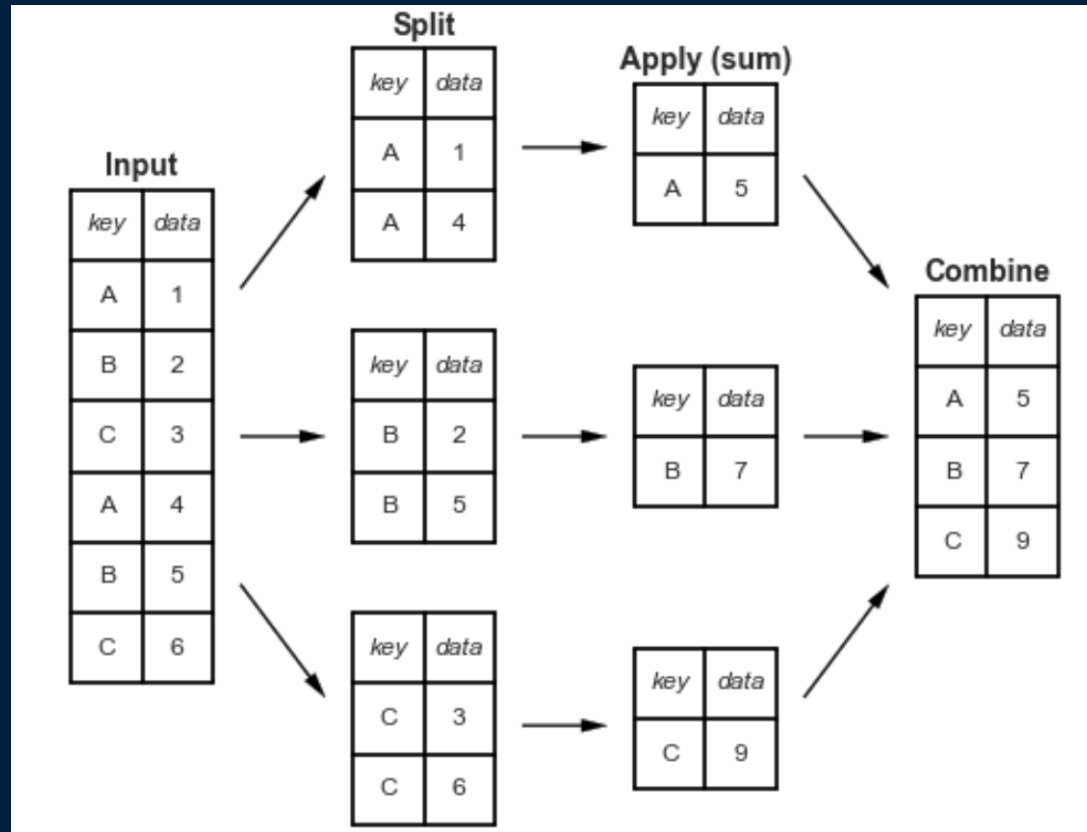
In `numpy` arrays, the aggregations output a single value

In `Pandas Series`, the aggregations output a single value

In `DataFrames`, the aggregations are operated in each column separately

You can set the `axis=1` or `axis='columns'` to operate the aggregation functions on each row

Deeper analytics using Groupby



Lazy evaluation

`groupby(column_name)` is the basic method:

`df.groupby('key')`

`<pandas.core.groupby.groupby.DataFrameGroupBy object at 0x11e7d9b38>`

Lazy evaluation: Nothing is done until an actual apply operation

`df.groupby('key').sum()`

	key	data
0	A	0
1	B	1
2	C	2
3	A	3
4	B	4
5	C	5

	key	data
	A	3
	B	5
	C	7

Try it yourself

Planet dataset from seaborn

It gives information on planets that astronomers have discovered around other stars (known as *extrasolar planets* or *exoplanets* for short).

Load the data:

```
import seaborn as sns
```

```
planets = sns.load_dataset('planets')
```

```
planets.shape
```

Take a moment to review the data in Pandas

Explore aggregation functions on the `planets` DataFrame or one column of the DataFrame

Quickly gain insights to a dataset

Use describe method on a DataFrame to get common aggregate functions: `planets.dropna().describe()`

	number	orbital_period	mass	distance	year
count	498.00000	498.000000	498.000000	498.000000	498.000000
mean	1.73494	835.778671	2.509320	52.068213	2007.377510
std	1.17572	1469.128259	3.636274	46.596041	4.167284
min	1.00000	1.328300	0.003600	1.350000	1989.000000
25%	1.00000	38.272250	0.212500	24.497500	2005.000000
50%	1.00000	357.000000	1.245000	39.940000	2009.000000
75%	2.00000	999.600000	2.867500	59.332500	2011.000000
max	6.00000	17337.500000	25.000000	354.000000	2014.000000

Discuss the results. What do you understand from this?

Groupby object

It is an abstraction of your data

Index a specific column with `groupby`: Referencing to a grouped Series by its column name

```
planets.groupby('method')['orbital_period'].median()  
# .count(), .sum(), etc...
```

```
method  
Astrometry                2  
Eclipse Timing Variations  9  
Imaging                   12  
Microlensing              7  
Orbital Brightness Modulation  3  
Pulsar Timing             5  
Pulsation Timing Variations  1  
Radial Velocity           553  
Transit                   397  
Transit Timing Variations  3  
Name: orbital_period, dtype: int64
```

Try It

Try an example from the *Python Data Science Handbook* by Jake VanderPlas

Go to the “Example: US States Data” section of the book:

<https://jakevdp.github.io/PythonDataScienceHandbook/03.07-merge-and-join.html>

Download the datasets and work through the example given in the book.



THE UNIVERSITY OF BRITISH COLUMBIA

