


# Lecture 6

Non-Parametric Density Estimation


# Motivating Example

We have some highly irregular data, how can we fit a distribution?



# Motivating Example

What if instead of one distribution, we fit several overlapping distributions?



# Mixture Models

- A **Gaussian Mixture Model** assumes the data is generated by  $K$  distributions, each with its own mean ( $\mu_k$ ), covariance ( $\Sigma_k$ ), and weight ( $\pi_k$ )
  - Non-Gaussian mixture models also exist, but we won't be covering those in this lecture
- To fit a Gaussian Mixture Model, we must do two things
  - Assign a probability of each data point being in each distribution
  - Fit each distribution using the methods we learned earlier

# Fitting GMMs

- The likelihood given by a Gaussian Mixture Model takes the equation

$$P(X_i) = \sum_{k=1}^K \pi_k P(X_i | Z_i = k)$$

- We can optimize this using the **expectation-maximization** algorithm
  - **Expectation:** Calculate the probability of a point being in each class
  - **Maximization:** Given these estimates, calculate using weighted maximum likelihood the parameters of each distribution.


# Mixture Model Demo

```
X1 = np.random.normal(10,5,size=300)
X2 = np.random.normal(20,5,size=100)
X3 = np.random.normal(50,10,size=250)
X = np.concatenate([X1,X2,X3]).reshape(-1,1)

from sklearn.mixture import GaussianMixture
gm = GaussianMixture(n_components=3)
gm.fit(X)

xr = np.linspace(X.min(), X.max(), 200).reshape(-1,1)
density = np.exp(gm.score_samples(xr))

print(gm.means_)
print(gm.covariances_)
print(gm.weights_)
```




# Non-Gaussian Mixture Models

- Non-Gaussian Mixture Models do exist (*e.g.* Poisson Mixture Models) but aren't commonly found in software packages
- You can fit any distribution mixture you want using the E-M algorithm (but I won't ask you to in this course)

# Another Example

I want to understand the distribution of how much customers spend in one purchase.

I have recorded the sale value for all my sales in the past week






# Why not use parameterized models?

- **Parameterized** models estimate based on a **finite** number of parameters
  - A normal distribution is parameterized by mean ( $\mu$ ) and standard deviation ( $\sigma$ )
- Parameterized models are also limited because the real world does not follow their assumptions
- **Non-parameterized** have huge, effectively **infinite**, numbers of parameters that allow them to be more expressive

# Estimating Density from Histograms

- We can directly convert a histogram into a probability density functions using the formula

$$f_H(x) = \frac{1}{nh} \sum_{i=1}^n 1(X_i \text{ in the same bin as } x)$$




# Estimating Density from Histograms

An alternative is the **Naïve Density Estimator**

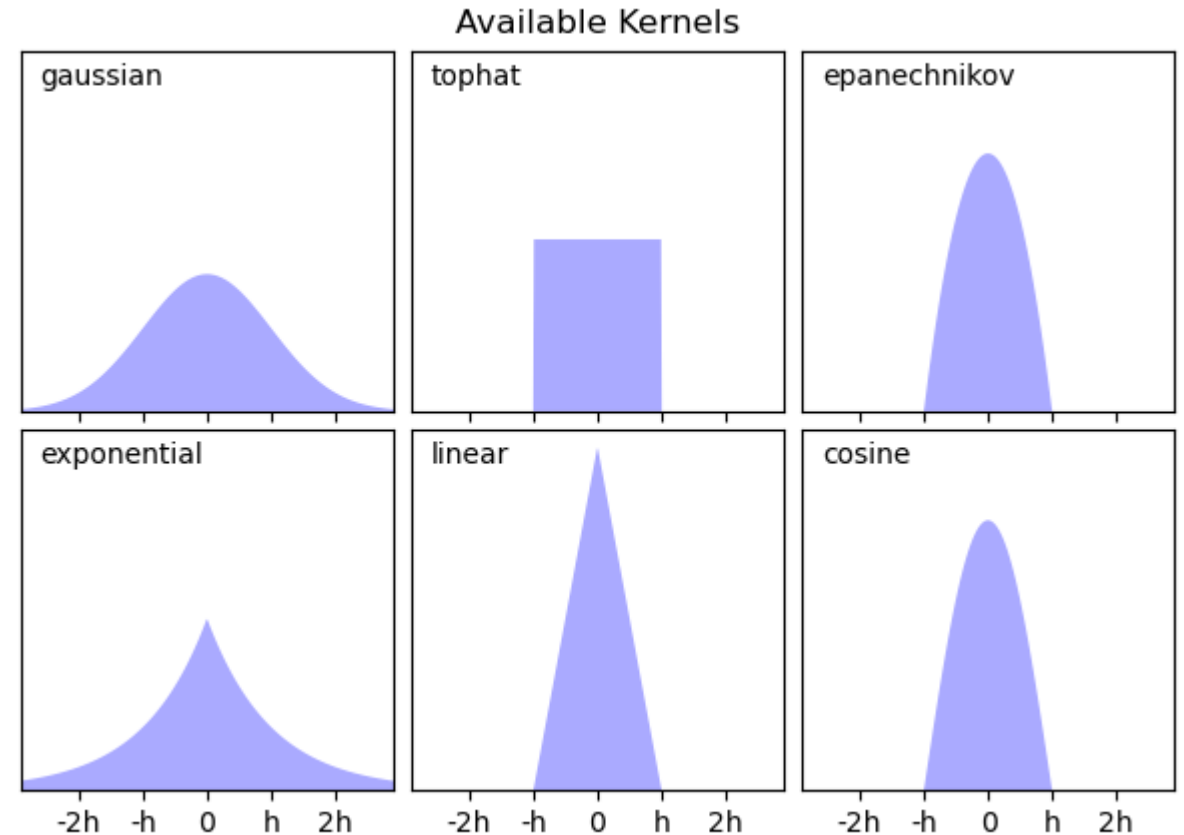
This essentially sets a histogram bin centered on  $x$

$$f_N(x) = \frac{1}{nh} \sum_{i=1}^n \frac{1((x-h) \leq x_i \leq (x+h))}{2}$$



# Introducing Kernels

- A **kernel** is a function that weights each data point based on its distance to the point of evaluation



# Estimating Density using Kernels

- Each kernel produces a probability distribution based on its data point


$$\hat{f}(x) = \frac{1}{nh} \sum_{i=1}^n K(z)$$

$$z = \frac{x - X_i}{h}$$

# Definition of non-parametric methods




- Let's look at an example using a **Gaussian kernel**

$$\hat{f}(x) = \frac{1}{nh} \sum_{i=1}^n \frac{1}{\sqrt{2\pi}} e^{\frac{1}{2}\left(\frac{x-X_i}{h}\right)^2}$$



# Effect of bandwidth

- **Bandwidth** is a scaling factor for the width of the kernel distributions



# Bandwidth selection

- The optimal bandwidth (*i.e.* bandwidth that minimizes the mean squared error) depends on the true probability distribution – which you don't know
- We can approximate this for Gaussian-esque data using **Silverman's rule of thumb**

$$\hat{h} = \left( \frac{4 * \hat{\sigma}^5}{3n} \right)^{\frac{1}{5}} \approx 1.06 * \hat{\sigma} * n^{-\frac{1}{5}}$$




# Finishing the example

```
from sklearn.neighbors import KernelDensity

sigma_hat = np.std(df['Sales'])
n = len(df['Sales'])
h_hat = 1.06 * sigma_hat * n**-0.2
kde = KernelDensity(kernel='gaussian', bandwidth=h_hat)
      .fit(df['Sales'])


xr = np.linspace(0, df['Sales'].max(), num=100)

# KernelDensity.score_samples returns log-likelihood,
# not density.
p = np.exp(kde.score_samples(xr))
```




# Limits of Silverman's Rule

- Silverman's rule of thumb assumes the data is near-normal. If your data differs from this, it tends to over-smooth



# Cross-Validation Bandwidth Selection


- Use grid search or another parameter optimization tool
  - Scikit-learn.model\_selection has good options here



# Another Example

Extending my previous study, I want to examine how many of each cookie flavour I sell.

I have recorded the flavour of the last 30 cookies sold.



# Density Estimation of Categorical Variables

- The simplest method is the relative frequency estimation

$$p(x) = \frac{1}{N} \sum_{i=1}^N \mathbb{1}_{X_i = x}$$

- In some cases, we may want to smooth this estimate using a smoothing kernel

$$p(x) = \frac{1}{N} \sum_{i=1}^N \begin{cases} 1 - \lambda & ; X_i = x \\ \frac{\lambda}{c - 1} & ; \text{otherwise} \end{cases}$$

# Bandwidth Selection for Categorical Variables

- For the estimator we selected, we can derive a plug-in estimate of the bandwidth

$$\lambda = \frac{c-1}{c} \left( 1 + \frac{n \sum_{x=1}^c \left[ \frac{1}{c} - p(x) \right]^2}{\sum_{x=1}^c p(x) [1 - p(x)]} \right)^{-1}$$

# Finding the Optimal Bandwidth

```
def find_bandwidth(data):  
    c = len(np.unique(data))  
    n = len(data)  
    _, counts = np.unique(data, return_counts=True)  
    probabilities = counts/sum(counts)  
  
    numerator = n*sum([(1/c + p)**2 for p in probabilities])  
    denominator = sum([p*(1-p) for p in probabilities])  
    return (c-1)/c * (1+numerator/denominator)**-1  
  
print(find_bandwidth(sales))  
>>> 0.3530903534154827
```

# Solving the Example

