

Lecture 9

Time Series


Basics of time series

- Random walks
- Autoregressive series
- Moving average series
- Autocorrelation and partial autocorrelation
- Differencing
- Seasonality

Example

My cookie factory wants to predict how many cookies will be sold in a day

We have recorded the number of sales every day for one year



Random Walks

- A random walk is a process in which each step is randomly determined
- The simplest random walk is an autoregressive random walk

$$X_t = X_{t-1} + \epsilon_t$$

Autoregressive Models

- An autoregressive (AR) is a generalization of random walks in which the value at time t depends on multiple preceding values

$$X_t = \phi_1 X_{t-1} + \phi_2 X_{t-2} + \cdots + \epsilon$$

- AR models typically do not skip lags. A model with three lags – AR(p) – can then be written as

$$X_t = \sum_{i=1}^p \phi_i X_{t-i}$$

Moving Average Models

- A moving average model has a value that depends on the last few error terms ϵ
- Moving average models – MA(q) – have equations of the form

$$X_t = \sum_{i=1}^q \epsilon_{t-i}$$

ARMA Models

- Models can have both autoregressive and moving average terms. These are called ARMA models and are parameterized by (p,q)

$$X_t = \sum_{i=1}^p \phi_i X_{t-1} + \sum_{j=1}^q \epsilon_{t-1} + \epsilon_t$$


- To select the lags p and q , we can look at the autocorrelation and partial autocorrelation

Autocorrelation & Partial Autocorrelation

```
fig = plt.figure(figsize=(12, 8))
ax1 = fig.add_subplot(211)

fig = sm.graphics.tsa.plot_acf(
    sales, lags=20, ax=ax1)
ax2 = fig.add_subplot(212)

fig =
sm.graphics.tsa.plot_pacf(
    sales, lags=20, ax=ax2)
```



Fitting Using Statsmodels

```
arma = sm.tsa.arima.ARIMA(data['Power'], order=(4,0,0)).fit()  
arma.summary()
```


```
=====
                        SARIMAX Results
=====
Dep. Variable:          y      No. Observations:      365
Model:                ARIMA(3, 0, 0)  Log Likelihood    -836.722
Date:                Tue, 12 Mar 2024  AIC              1683.444
Time:                 13:40:58  BIC              1702.943
Sample:              0      HQIC              1691.193
                        - 365
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	42.1426	2.929	14.387	0.000	36.401	47.884
ar.L1	0.7633	0.053	14.356	0.000	0.659	0.867
ar.L2	0.3338	0.065	5.161	0.000	0.207	0.461
ar.L3	-0.1388	0.052	-2.657	0.008	-0.241	-0.036
sigma2	5.6971	0.453	12.564	0.000	4.808	6.586


```
=====
Ljung-Box (L1) (Q):      0.01  Jarque-Bera (JB):      0.65
Prob(Q):                0.91  Prob(JB):          0.72
Heteroskedasticity (H):  1.04  Skew:              0.04
Prob(H) (two-sided):    0.85  Kurtosis:          2.81
=====
```

Using the results



Short-Term Forecasting

Predict values over the next few time periods




Representative Scenarios

Create sample time series following the same patterns

Stationary and Non-Stationary Series

- ARMA models assume the time series is stationary, meaning that the distribution of values does not change as time passes
- What happens if we do not have a stationary time series?



Non-Stationary Time Series



- We can check if a time series is stationary using the Augmented Dickey-Fuller (ADF or ADFuller) test
 - We won't go into details on how this is calculated, but it ends with a hypothesis test that gives a p-value. Lower p-value means more stationary

```
sm.tsa.adfuller(ts)
>>> (0.18680600896643582, # Test statistic
      0.9715117981353568, # P-Value
      0, 99, # Lags used, observations
      {'1%': -3.498198082189098, # critical values
       '5%': -2.891208211860468,
       '10%': -2.5825959973472097},
      526.2316331583049) # best IC score
```

Non-Stationary Time Series


- For non-stationary time series, we can introduce one or more differences to the ARMA model – creating an ARIMA model
 - An ARIMA($p,1,q$) model is equivalent to an ARMA(p,q) model on the differences of the time series
- You can estimate terms for p and q using ACF/PACF plots of the **difference** in the time series

Non-Stationary Time Series



Seasonality


- Seasonality is when time series data follows a recurring pattern over some interval
- We can decompose seasonal data into seasonal, trend, and residual terms
 - Seasonal – repeating pattern
 - Trend – long-term average
 - Residual – leftovers



STL Decomposition

- **Seasonal-Trend** Decomposition by **Loess** uses smoothing estimates to generate the component time series

```
from statsmodels.tsa.seasonal import STL
stl = STL(df['Power'], period=24)
result = stl.fit()
```



STL Decomposition

- The residuals of an STL decomposition can be treated as an ARIMA series

```
from statsmodels.tsa.seasonal import STL
stl = STL(df['Power'], period=24)
result = stl.fit()

res = sm.tsa.ARIMA(result.resid,
                    order=(1,0,1)).fit()
```

